

Milestone 3 Report

0. Baseline:

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	<i>0.189199 ms</i>	<i>0.515635 ms</i>	<i>0m4.078s</i>	<i>0.86</i>
1000	<i>1.20908 ms</i>	<i>5.04132 ms</i>	<i>0m36.842s</i>	<i>0.886</i>
10000	<i>11.5888 ms</i>	<i>53.6658 ms</i>	<i>6m10.192s</i>	<i>0.8714</i>

1. Optimization Number #0: Streams

- a. How does this optimization work in theory? Expected behavior?

This optimization (streams) allows for concurrent execution of kernel computations and memory transfers. By using streams, data transfers can be overlapped during kernel executions which can reduce op times. The streams can operate in separate queues at once. Op times may improve with streams as data transfers overlap with kernel execution reducing memory latency, but it may not be significant.

- b. How did you implement your code? Explain thoroughly and show code snippets.

Multiple streams are created to manage parallel execution of data transfers and kernel launches using `cudaStreamCreate`. The input data, mask, and outputs are managed in segments to facilitate batch processing. Each stream handles a segment of the batch, and this is done by using offsets for the input and output pointers. In this code, there are 4 streams that divide the batch up. Asynchronous memory transfers

(cudaMemcpyAsync) and kernel executions are scheduled on different streams. After all operations are scheduled, the streams are destroyed.

```
cudaStream_t stream[numStreams];
for (int i = 0; i < numStreams; ++i){
    cudaStreamCreate(&stream[i]);
}

for(int i = 0; i < numStreams; ++i){
    int start = i * segsize;
    int currentBatchSize = min(segSize, Batch - start); //batch size adjusted for in case of smaller batch

    if (currentBatchSize > 0) {
        cudaMemcpyAsync(*device_input_ptr + start * segSizeIn, host_input + start * segSizeIn, currentBatchSize * segSizeIn * sizeof(float), cudaMemcpyHostToDevice, stream[i]);
        cudaMemcpyAsync(*device_mask_ptr, host_mask, maskSize * sizeof(float), cudaMemcpyHostToDevice, stream[i]);

        conv_forward_kernel<<<dim3Grid, dim3Block, 0, stream[i]>>>(*device_output_ptr + start * segSizeOut, *device_input_ptr + start * segSizeIn, *device_mask_ptr, Batch, Map_out, Channel, Height, Width, K);
        cudaMemcpyAsync((float*)host_output + start * segSizeOut, *device_output_ptr + start * segSizeOut, currentBatchSize * segSizeOut * sizeof(float), cudaMemcpyDeviceToHost, stream[i]);
    }
}

for (int i = 0; i < numStreams; ++i)
    cudaStreamDestroy(stream[i]);
}
```

Fig. – Code snippet of stream usage for copying memory and kernel launches

- c. Did the performance match your expectation? Show your analysis results using profiling tools.

*Layer Times are recorded here due to streams being written in conv_forward_gpu_prolog

Batch Size	Layer Time 1	Layer Time 2	Total Execution Time	Accuracy
100	2.89909 ms	2.96913 ms	0m4.062s	0.86
1000	20.6266 ms	18.9936 ms	0m37.565s	0.886
10000	177.632 ms	181.944 ms	6m5.448s	0.8714

While it is difficult to identify if streams improve the code, we can identify op times through looking at the total kernal time in the nsys profiling below. For a batch of 10,000, the first layer the op time is 52,964,741ns and for the second layer, the op time is 11,702,168ns. Comparing these results closely to the op times for the baseline, it can be noted that there is no improvement. The lack of improvement can be explained by the baseline code being the same, and the only time is saved on copying memory, which is negligible in this case. The noticeable improvement in the nsys profiling is the large reduction in cudaapisum where cudaDeviceSynchronize nearly consumes no time. Compared to other profiles, cudaDeviceSynchronize was always

in the top three of largest times in `cudaapisum`. Also, in `thegpumemtimesum`. The overall time is lower compared to other optimizations.

Nsight compute shows that the GPU throughput is very high and balanced. This is expected as the forward kernel code does not change. The code can improve with its memory utilization as the L2 load and store access pattern is not optimal show in the Nsight memory workload analysis below. Making sure that memory accesses take advantage of coalescing can improve performance.

[5/8] Executing 'cudaapisum' stats report									
Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name	
69.8	353,961,951	30	11,798,731.7	4,181,818.0	5,380	33,683,445	14,105,672.9	cudaMemcpyAsync	
29.9	151,531,253	8	18,941,406.6	127,683.5	8,646	150,686,833	53,233,246.7	cudaMalloc	
0.1	573,754	2	286,877.0	286,877.0	35,306	538,448	355,775.1	cudaMemcpy	
0.1	358,182	14	25,584.4	10,339.5	7,784	95,459	29,230.6	cudaLaunchKernel	
0.0	134,773	10	13,477.3	4,864.5	3,847	76,293	22,548.1	cudaStreamDestroy	
0.0	133,130	8	16,641.3	1,999.0	501	189,144	37,570.7	cudaFree	
0.0	53,272	10	5,327.2	3,085.5	1,723	15,998	5,547.9	cudaStreamCreate	
0.0	33,783	6	5,630.5	5,475.5	4,689	7,483	1,014.6	cudaDeviceSynchronize	
0.0	1,042	1	1,042.0	1,042.0	1,042	1,042	0.0	cuModuleGetLoadingMode	

[6/8] Executing 'gpu kernsum' stats report									
Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ
81.9	52,964,741	5	10,592,948.2	10,581,665.0	10,578,017	10,612,545	17,285.9	2000	16 9 16 16 1
18.1	11,782,168	5	2,340,433.6	2,340,190.0	2,339,774	2,341,534	672.9	2000	4 25 16 16 1
0.0	4,688	2	2,304.0	2,304.0	2,272	2,336	45.3	1 1 1 1 1	prefn_marker_kernel()
0.0	4,543	2	2,271.5	2,271.5	2,239	2,304	46.0	1 1 1 1 1	do_not_remove_this_kernel()

[7/8] Executing 'gpumemtimesum' stats report									
Time (%)	Total Time (ns)	Count	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Operation	
85.7	245,492,370	10	24,549,237.0	25,875,081.0	19,744,449	29,004,430	4,052,453.9	[CUDA memcpy DtoH]	
14.3	40,838,899	22	1,856,313.6	2,160.0	1,215	4,891,195	2,184,494.6	[CUDA memcpy HtoD]	

[8/8] Executing 'gpumemsizesum' stats report							
Total (MB)	Count	Avg (MB)	Med (MB)	Min (MB)	Max (MB)	StdDev (MB)	Operation
1,763.840	10	176.384	176.384	147.968	204.800	29.953	[CUDA memcpy DtoH]
551.907	22	25.087	0.013	0.000	59.168	28.256	[CUDA memcpy HtoD]

Fig. – nsys result 5-8 for Streams

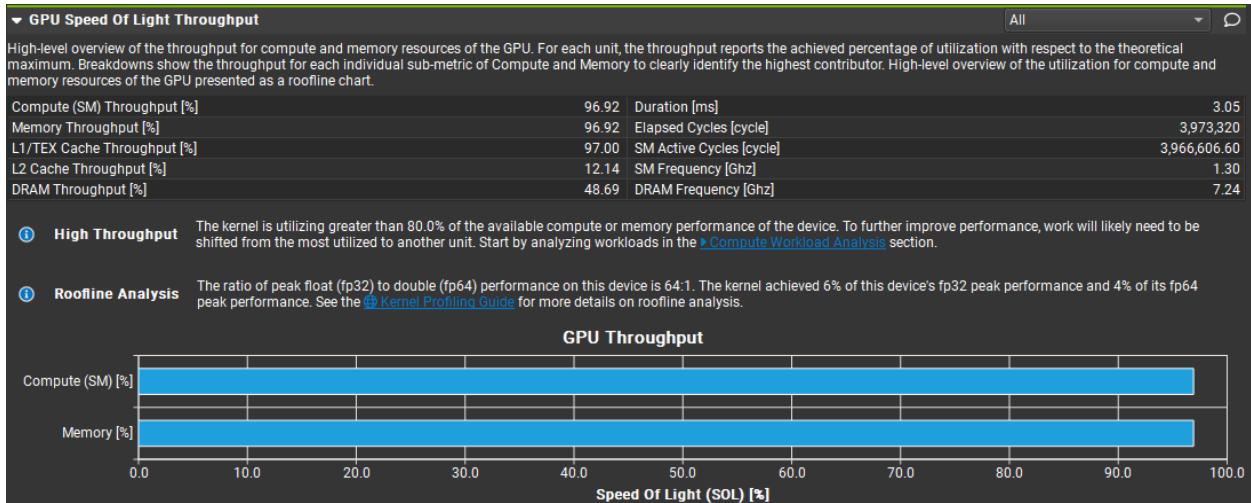


Fig. – Streams GPU SOL Throughput

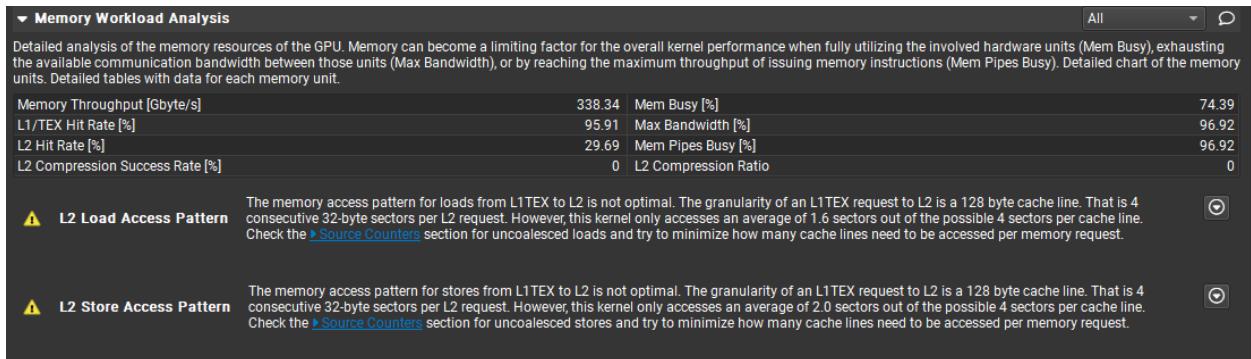


Fig. – Memory workload analysis for Streams

d. Does this optimization synergize with any other optimizations? How and why?

Using constant memory with streams optimizes performance by reducing memory bandwidth and improving data access speed. Constant memory provides fast access to data across multiple streams, enhancing efficiency and throughput in parallel kernel executions. Using constant memory can eliminate the need for multiple copies of the mask, streamlining memory usage. Additionally, employing streams can further optimize performance by reducing data transfer times for optimizations that may need multiple kernel launches such as input matrix unrolling & tiled multiplication using shared memory.

e. List your references used while implementing this technique. (you must mention textbook pages at the minimum)

- i. Slide Deck 20: GPU Data Transfer
- ii. Cuda C++ Programming Guide, Release 12.3 – p. 48

2. Optimization Number #1: Tiled (Shared Memory) Convolution

- a. How does this optimization work in theory? Expected behavior?

Tiling, or using shared memory for convolution, is an optimization technique to reduce global memory accesses by loading chunks of the input data and kernel into faster shared memory. Shared memory is much faster than global memory, but its size is limited, so it's used to store smaller tiles of the input and kernel that are then used to compute a portion of the output. By loading data into shared memory once and reusing it for multiple computations, we minimize the number of slow global memory accesses.

With the use of shared memory, reduced global memory access and increased throughput can be expected.

- b. How did you implement your code? Explain thoroughly and show code snippets.

The implementation leverages shared memory to store input tiles and convolution masks, reducing the number of global memory accesses. Below is an image of the code where the tiles in shared memory are loading the kernel and input. After loading the necessary input elements, the convolution is performed and written back to global memory.

```

int n = blockIdx.x;
int m = blockIdx.y;
int w_tx = threadIdx.x;
int h_ty = threadIdx.y;
int hBase = (blockIdx.z/W_grid) * TILE_WIDTH;
int wBase = (blockIdx.z % W_grid) * TILE_WIDTH;
int h = hBase + w_tx;
int w = wBase + h_ty;

float acc = 0.0;
for(int c = 0; c < Channel; c++){
    //loading shared mask
    if(w_tx < K && h_ty < K){
        W_shared[h_ty * K + w_tx] = mask_4d(m, c, h_ty, w_tx);
    }

    __syncthreads();
    //loading shared input
    for(int i = h; i < hBase + X_tile_width; i += TILE_WIDTH){
        for (int j = w; j < wBase + X_tile_width; j += TILE_WIDTH){
            if(i < Height && j < Width){
                X_shared[(i - hBase) * X_tile_width + (j - wBase)] = in_4d(n, c, i, j);
            }
        }
    }
}

```

Fig. – Code snippet of loading shared tiles with inputs

- c. Did the performance match your expectation? Show your analysis results using profiling tools.

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	<i>0.191784 ms</i>	<i>0.707698 ms</i>	<i>0m4.096s</i>	<i>0.86</i>
1000	<i>1.33759 ms</i>	<i>6.90021 ms</i>	<i>0m37.074s</i>	<i>0.886</i>
10000	<i>13.0311 ms</i>	<i>68.7664 ms</i>	<i>6m3.390s</i>	<i>0.8714</i>

As seen by the op times, this optimization increases the run time. This is the opposite of the expected behavior. The overall run time increases as time is spent copying over the data to shared memory and bank conflicts. Looking at the nsys profiling results below, it can be noticed that the `conv_forward_kernel` runs for a greater amount of time, compared to successful optimizations.

[6/8] Executing 'gpukernsum' stats report												Name
Time (%)	Total	Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ		Name
84.2	68,841,799		1	68,841,799.0	68,841,799.0	68,841,799	68,841,799	0.0	10000	16	9	conv_forward_kernel(float *, const float *, const float *, int, int)
15.8	12,935,557		1	12,935,557.0	12,935,557.0	12,935,557	12,935,557	0.0	10000	4	25	conv_forward_kernel(float *, const float *, const float *, int, int)
0.0	4,768		2	2,384.0	2,384.0	2,368	2,400	22.6	1	1	1	do_not_remove_this_kernel()
0.0	4,705		2	2,352.5	2,352.5	2,305	2,400	67.2	1	1	1	prefn_marker_kernel()

[7/8] Executing 'gpumemtimesum' stats report											
Time (%)	Total	Time (ns)	Count	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Operation		
85.6	248,216,895		2	124,198,447.5	124,198,447.5	101,686,378	146,530,517	31,709,594.8	[CUDA memcpy DtoH]		
14.4	41,762,183		6	6,960,350.5	1,840.5	1,536	22,042,003	10,805,422.9	[CUDA memcpy HtoD]		

[8/8] Executing 'gpumemsizesum' stats report							
Total (MB)	Count	Avg (MB)	Med (MB)	Min (MB)	Max (MB)	StdDev (MB)	Operation
1,763.848	2	881.920	881.920	739.848	1,024.000	200.931	[CUDA memcpy DtoH]

Fig. – nsys results 6-8 for Tiled (Shared Memory) Convolution

When analyzing in Nsight Compute, GPU throughput slightly decreases from the baseline as seen below.

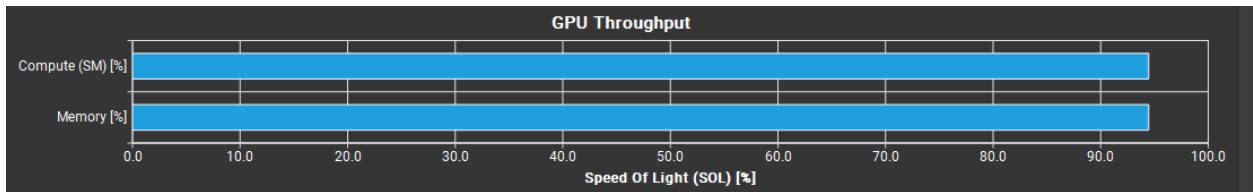


Fig. – Tiled (Shared Memory) Convolution GPU SOL Throughput

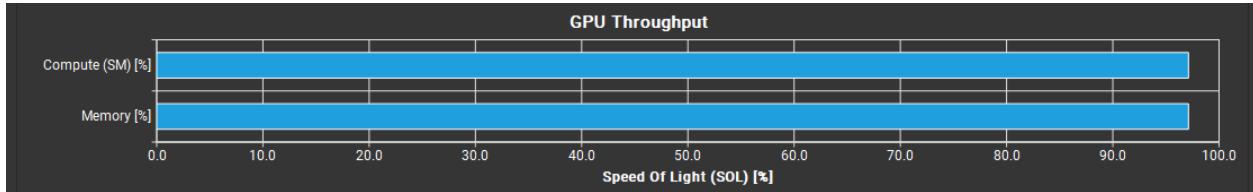


Fig. – Baseline GPU SOL Throughput

Another issue can be found in the L1TEX global load/store access pattern and the L2 load/store access pattern as seen below. Looking at the prompts next to the yellow icons, is the explanation for why there are problems with the access patterns for the L1TEX and L2. To improve functionality, global memory accesses should accommodate to the number of bytes in a burst. Additionally, in the figure below showing data on shared memory and L1/TEX Cache, it is visible that there are 7,224,461 bank conflicts and misses to the L2 cache. By optimizing global memory accesses, these values can be reduced. With the large count of bank conflicts within the shared memory, run time slows down which eliminates the benefit of using shared memory. Improving memory access patterns in this optimization could possibly lead to a benefit.

Detailed analysis of the memory resources of the GPU. Memory can become a limiting factor for the overall kernel performance when fully utilizing the involved hardware units (Mem Busy), exhausting the available communication bandwidth between those units (Max Bandwidth), or by reaching the maximum throughput of issuing memory instructions (Mem Pipes Busy). Detailed chart of the memory units. Detailed tables with data for each memory unit.

Memory Throughput [Gbyte/s]	334.63	Mem Busy [%]	55.81
L1/TEX Hit Rate [%]	65.49	Max Bandwidth [%]	94.46
L2 Hit Rate [%]	47.69	Mem Pipes Busy [%]	94.46
L2 Compression Success Rate [%]	0	L2 Compression Ratio	0

⚠ L1TEX Global Load Access Pattern The memory access pattern for global loads in L1TEX might not be optimal. On average, this kernel accesses 4.0 bytes per thread per memory request; but the address pattern, possibly caused by the stride between threads, results in 9.5 sectors per request, or $9.5 \times 32 = 302.4$ bytes of cache data transfers per request. The optimal thread address pattern for 4.0 byte accesses would result in $4.0 \times 32 = 128.0$ bytes of cache data transfers per request, to maximize L1TEX cache performance. Check the [Source Counters](#) section for uncoalesced global loads.

⚠ L1TEX Global Store Access Pattern The memory access pattern for global stores in L1TEX might not be optimal. On average, this kernel accesses 4.0 bytes per thread per memory request; but the address pattern, possibly caused by the stride between threads, results in 16.0 sectors per request, or $16.0 \times 32 = 512.0$ bytes of cache data transfers per request. The optimal thread address pattern for 4.0 byte accesses would result in $4.0 \times 32 = 128.0$ bytes of cache data transfers per request, to maximize L1TEX cache performance. Check the [Source Counters](#) section for uncoalesced global stores.

⚠ L2 Store Access Pattern The memory access pattern for stores from L1TEX to L2 is not optimal. The granularity of an L1TEX request to L2 is a 128 byte cache line. That is 4 consecutive 32-byte sectors per L2 request. However, this kernel only accesses an average of 1.2 sectors out of the possible 4 sectors per cache line. Check the [Source Counters](#) section for uncoalesced stores and try to minimize how many cache lines need to be accessed per memory request.

⚠ L2 Load Access Pattern The memory access pattern for loads from L1TEX to L2 is not optimal. The granularity of an L1TEX request to L2 is a 128 byte cache line. That is 4 consecutive 32-byte sectors per L2 request. However, this kernel only accesses an average of 1.1 sectors out of the possible 4 sectors per cache line. Check the [Source Counters](#) section for uncoalesced loads and try to minimize how many cache lines need to be accessed per memory request.

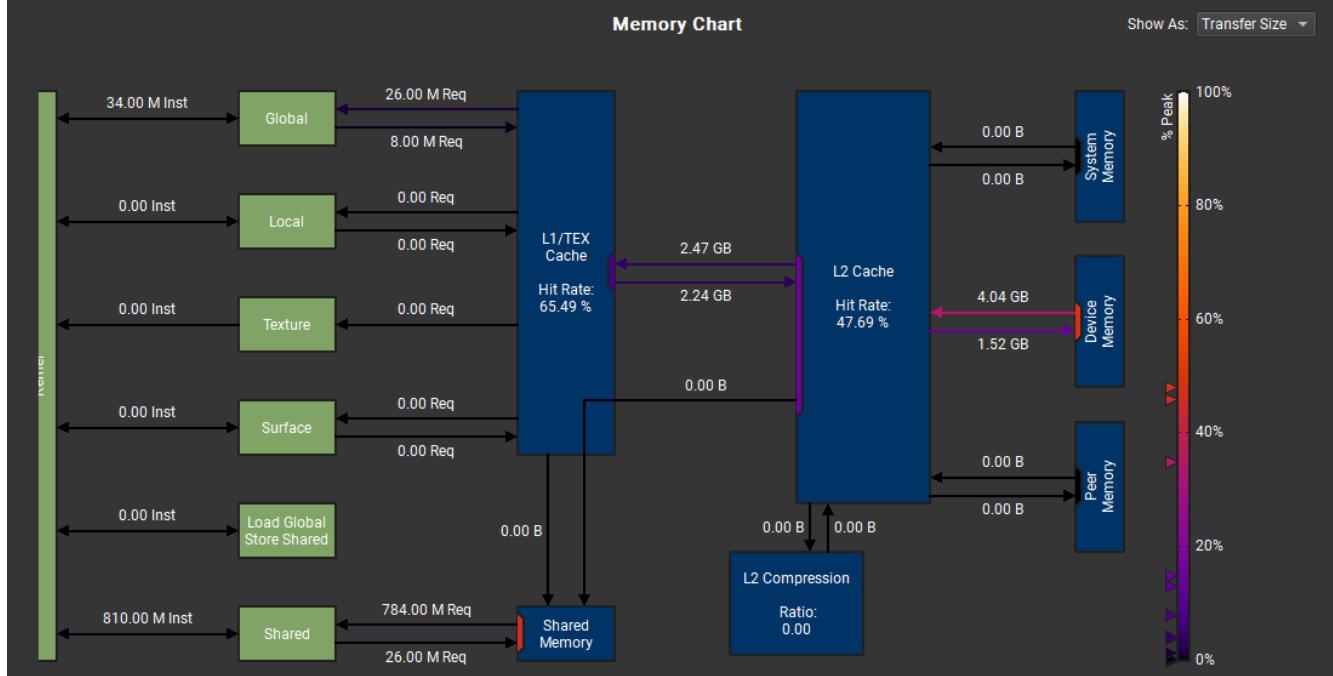


Fig. – Nsight Compute Memory Chart for Tiled (Shared Memory) Convolution

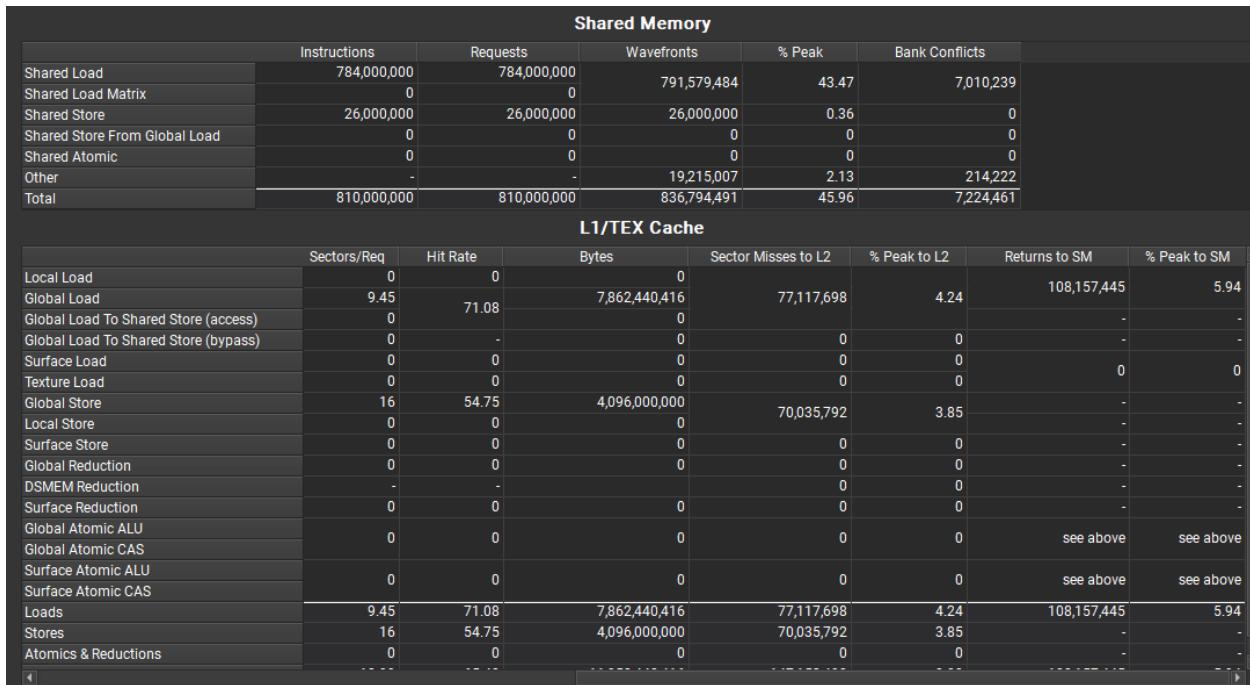


Fig. – Nsight compute data chart for shared memory and L1/TEX cache

d. Does this optimization synergize with any other optimizations? How and why?

Yes, this optimization can synergize with constant memory usage by storing frequently accessed kernel data in constant memory, we can further reduce global memory accesses. To improve the global memory access pattern, focusing on memory coalescing can reduce the flaws of this optimization. Additionally, loop unrolling within the shared memory kernel can reduce the overhead of loop control and increase instruction-level parallelism.

e. List your references used while implementing this technique. (you must mention textbook pages at the minimum)

i. Course Textbook p. 358

3. Optimization Number #2: Input Matrix Unrolling & Tiled Multiplication using Shared Memory

a. How does this optimization work in theory? Expected behavior?

The optimization combines two techniques, Input Matrix Unrolling and Tiled Multiplication using Shared Memory. Input matrix unrolling involves transforming the input matrix into a different representation to optimize memory access patterns. Tiled matrix multiplication is like the tiled convolution optimization, this approach loads tiles of the input matrix and kernel into shared memory to minimize global memory accesses and utilize the faster shared memory for computations. Since the matrix is unrolled explicitly (in a separate kernel), the run time is expected to slow down because the data is being replicated and transferred redundantly.

b. How did you implement your code? Explain thoroughly and show code snippets.

The implementation includes input matrix unrolling and tiled matrix multiplication using shared memory in two separate kernels. The unrolling kernel transforms the input matrix into a more optimized format for subsequent processing. This is done by rearranging the input data into a linear array that facilitates coalesced memory accesses and efficient computation in the next stage. The tiled matrix multiplication kernel performs the matrix multiplication using tiles of the input and kernel data loaded into shared memory. Lastly, the kernels are launched in “mini batches” using a loop, since running a large batch (above ~5000 images) leads to code failure which could be a result of global memory limitations. Snippets of the code for each kernel and kernel launches can be seen below.

```
//unroll kernel
__global__ void unroll_kernel(int b, int Channel, int Height, int Width, int K, const float* input, float* X_unroll) {
    int t = blockIdx.x * blockDim.x + threadIdx.x;
    int Height_out = Height - K + 1;
    int Width_out = Width - K + 1;
    int W_unroll = Height_out * Width_out;
    int H_unroll = Channel * K * K;

    int by = blockIdx.y;

#define X_unroll_output(i2, i1, i0) X_unroll[(i2) * (Width_out*Height_out*Channel*K*K) + (i1) * (Width_out * Height_out) + (i0)]
#define in_4d(i3, i2, i1, i0) input[(i3) * (Channel * Height * Width) + (i2) * (Height * Width) + (i1) * (Width) + i0]

    if (t < Channel * W_unroll) {
        int c = t / W_unroll;
        int s = t % W_unroll;
        int h_out = s / Width_out;
        int w_out = s % Width_out;
        int w_unroll = h_out * Width_out + w_out;
        int w_base = c * K * K;

        for (int p = 0; p < K; p++) {
            for (int q = 0; q < K; q++) {
                int h_unroll = w_base + p * K + q;
                X_unroll_output(by, h_unroll, w_unroll) = in_4d(by + b, c, h_out + p, w_out + q);
            }
        }
    }
}
```

Fig. – Matrix unrolling kernel

```
//matrix mult kernel
__global__ void matrixMultiplyShared(int b, const float *A, float *B, float *C,
                                     int numARows, int numAColumns,
                                     int numBRows, int numBColumns,
                                     int numCRows, int numCColumns, int Batch) {
    //@@ Insert code to implement matrix multiplication here
    //@@ You have to use shared memory for this MP
    __shared__ float subTileA[TILE_WIDTH][TILE_WIDTH];
    __shared__ float subTileB[TILE_WIDTH][TILE_WIDTH];

    int bx = blockIdx.x;
    int tx = threadIdx.x;
    int by = blockIdx.y;
    int ty = threadIdx.y;
    int bz = blockIdx.z;

    int Row = by * TILE_WIDTH + ty;
    int Col = bx * TILE_WIDTH + tx;
    float Cvalue = 0;

    for(int m = 0; m < ceil(numAColumns/(1.0 * TILE_WIDTH)) ; ++m){
        if(Row < numARows && (m * TILE_WIDTH + tx) < numAColumns){
            subTileA[ty][tx] = A[Row * numAColumns + (m * TILE_WIDTH + tx)];
        }
        else{
            subTileA[ty][tx] = 0;
        }

        if((m * TILE_WIDTH + ty) < numBRows && Col < numBColumns){
            subTileB[ty][tx] = B[(bz) * numBRows * numBColumns + (m * TILE_WIDTH + ty) * numBColumns + Col];
        }
        else{
            subTileB[ty][tx] = 0;
        }

        __syncthreads();

        for(int k = 0; k < TILE_WIDTH; ++k){
            Cvalue += subTileA[ty][k] * subTileB[k][tx];
        }
        __syncthreads();
    }

    if(Row < numCRows && Col < numCColumns && (bz + b) < Batch){
        C[(bz + b) * numCRows * numCColumns + Row * numCColumns + Col] = Cvalue;
    }
}
```

Fig. – Code snippet of tiled matrix multiplication kernel

```
for(int b = 0; b < Batch; b += mini_batch){
    unroll_kernel<<<dimGridUnroll, dimBlockUnroll>>>(b, Channel, Height, Width, K, device_input, X_unroll);
    cudaDeviceSynchronize();
    matrixMultiplyShared<<<dimGridMult, dimBlockMult>>>(b, device_mask, X_unroll, device_output, Map_out, H_unroll, H_unroll, W_unroll, Map_out, W_unroll, Batch);
    cudaDeviceSynchronize();
```

Fig. - Code snippet of “mini batch” launches

- c. Did the performance match your expectation? Show your analysis results using profiling tools.

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
------------	-----------	-----------	----------------------	----------

100	<i>1.17789 ms</i>	<i>0.961721 ms</i>	<i>0m4.086s</i>	<i>0.86</i>
1000	<i>8.52241 ms</i>	<i>5.80802 ms</i>	<i>0m37.380s</i>	<i>0.886</i>
10000	<i>82.477 ms</i>	<i>56.9284 ms</i>	<i>6m4.185s</i>	<i>0.8714</i>

Compared to the baseline, this optimization did not improve op times. This is due to the replication of the data. Having separate kernels requires the unrolled matrix to be computed separately and then transferred to the matrix multiplication. Such redundancies are nonoptimal for speed. Based on the nsys results below, there are 10 instances of each kernel launch, and each isn't very fast. In cudaapisum, it is also visible that there are 46 instances of cudaDeviceSynchronize since it needs to be executed after every kernel. Such factors slow the program down.

[5/8] Executing 'cudaapisum' stats report										
Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name		
51.1	321,532,491	8	40,191,561.4	10,570,340.0	16,291	163,398,324	63,125,807.9	cudaMemcpy		
25.3	159,366,709	10	15,936,678.9	148,482.0	88,911	157,625,495	49,784,605.4	cudaMalloc		
21.9	137,551,549	46	2,990,251.1	2,675,915.5	3,236	5,558,618	1,647,649.1	cudaDeviceSynchronize		
1.6	10,291,455	10	1,029,145.5	602,436.0	115,035	2,740,021	956,669.7	cudaFree		
0.1	526,525	44	11,966.5	5,255.0	4,088	107,572	21,419.3	cudaLaunchKernel		
0.0	1,042	1	1,042.0	1,042.0	1,042	1,042	0.0	cuModuleGetLoadingMode		

[6/8] Executing 'gpkernrsum' stats report										
Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ	
40.4	55,529,143	10	5,552,914.3	5,552,920.5	5,552,201	5,553,481	316.0	400 1 1000	16 16 1	matrixMultiplyShared(int, c
21.4	29,406,449	10	2,940,644.9	2,940,603.0	2,940,172	2,941,259	331.0	73 1 1000	16 16 1	matrixMultiplyShared(int, c
19.4	26,647,528	10	2,664,752.8	2,665,322.0	2,650,410	2,675,083	8,341.8	19 1000 1	256 1 1	unroll_kernel(int, int, int
18.7	25,687,754	10	2,568,775.4	2,568,260.0	2,567,396	2,572,196	1,362.8	25 1000 1	256 1 1	unroll_kernel(int, int, int
0.0	4,864	2	2,432.0	2,432.0	2,400	2,464	45.3	1 1 1	1 1 1	do_not_remove_this_kernel()
0.0	4,608	2	2,304.0	2,304.0	2,304	2,304	0.0	1 1 1	1 1 1	prefn_marker_kernel()

[7/8] Executing 'gpumemtimesum' stats report									
Time (%)	Total Time (ns)	Count	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Operation	
86.6	277,261,553	2	138,630,776.5	138,630,776.5	114,381,747	162,879,806	34,293,306.4	[CUDA memcpy DtoH]	
13.4	42,891,634	6	7,148,605.7	1,824.0	1,536	21,980,642	11,077,200.4	[CUDA memcpy HtoD]	

[8/8] Executing 'gpumemsizesum' stats report							
Total (MB)	Count	Avg (MB)	Med (MB)	Min (MB)	Max (MB)	StdDev (MB)	Operation
1,763.840	2	881.920	881.920	739.840	1,024.000	200.931	[CUDA memcpy DtoH]
551.853	6	91.976	0.007	0.000	295.840	143.039	[CUDA memcpy HtoD]

Fig. – nsys results 5-8 for Input Matrix Unrolling & Tiled Multiplication using Shared Memory

Looking at the Nsight Compute output below, it is realized that the GPU throughput for the matrix unrolling is nonoptimal. The compute percentage is very low while the memory usage is very high, underutilizing compute resources. In the figure containing the warp statistics for the matrix unroll kernel, the errors can be seen next to the yellow icon. In the unroll kernel, time is being wasted waiting to

receive data resulting in stalls and under utilization of the possible warp output shown in the warp statistics.

On the other hand, GPU throughput for the matrix multiplication kernel is much better and is balanced. The matrix multiplication kernel can improve in its L2 hit rate as it is at 7.76%. Loads and store from the L1/TEX to the L2 are not optimal as seen in the figure containing the memory workload analysis. The matrix multiplication could benefit from better memory access patterns.

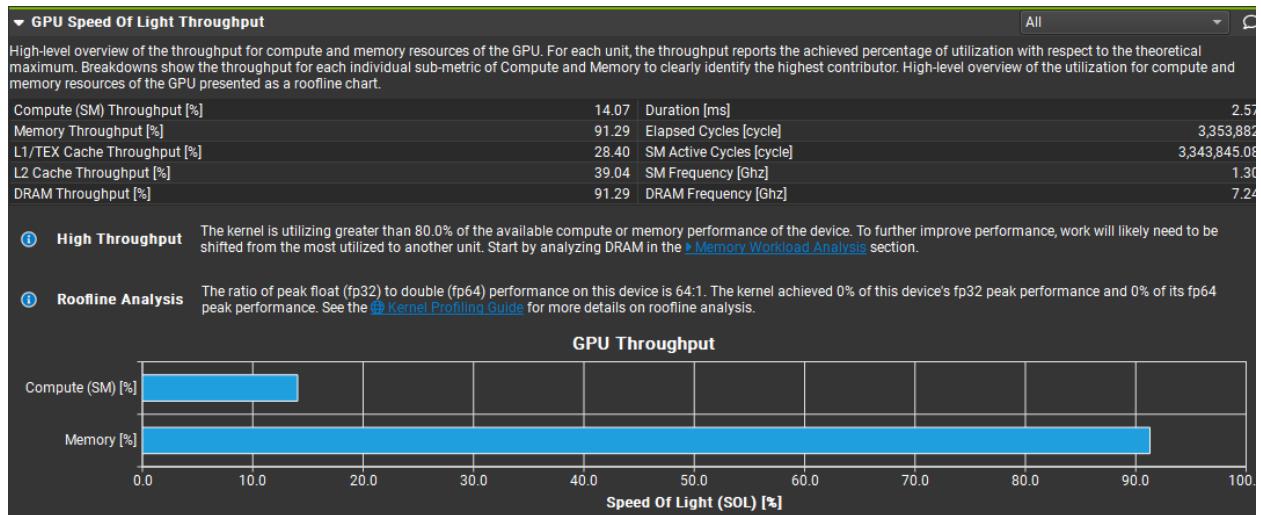


Fig. – Matrix Unroll Kernel GPU SOL Throughput

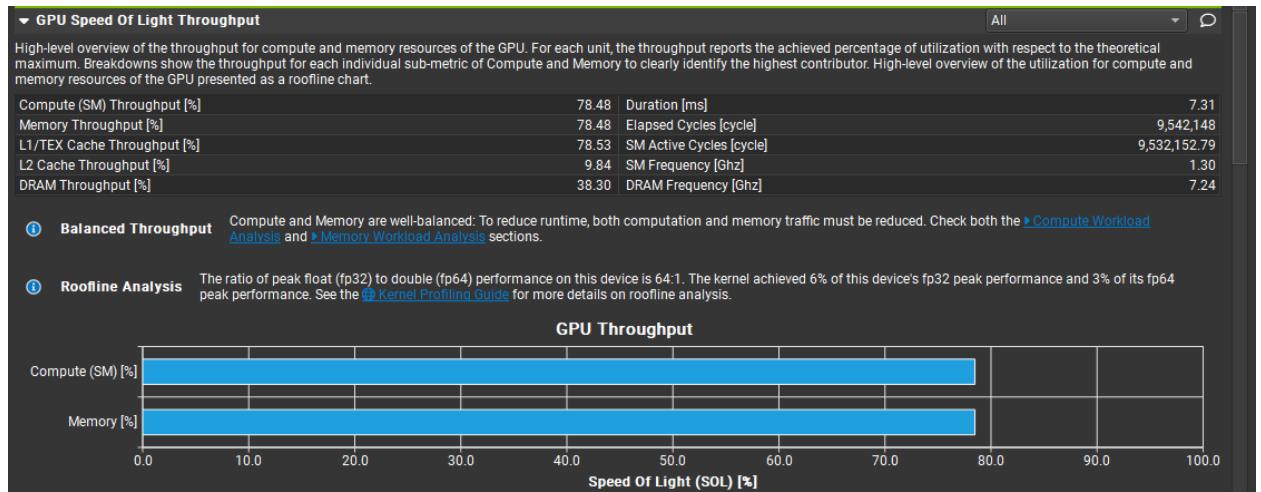


Fig. – Tiled Matrix Multiply GPU SOL Throughput

Scheduler Statistics

Summary of the activity of the schedulers issuing instructions. Each scheduler maintains a pool of warps that it can issue instructions for. The upper bound of warps in the pool (Theoretical Warps) is limited by the launch configuration. On every cycle each scheduler checks the state of the allocated warps in the pool (Active Warps). Active warps that are not stalled (Eligible Warps) are ready to issue their next instruction. From the set of eligible warps the scheduler selects a single warp from which to issue one or more instructions (Issued Warp). On cycles with no eligible warps, the issue slot is skipped and no instruction is issued. Having many skipped issue slots indicates poor latency hiding.

Active Warps Per Scheduler [warp]	10.22	No Eligible [%]	90.72
Eligible Warps Per Scheduler [warp]	0.15	One or More Eligible [%]	9.28
Issued Warp Per Scheduler	0.09		

Issue Slot Utilization

Every scheduler is capable of issuing one instruction per cycle, but for this kernel each scheduler only issues an instruction every 10.8 cycles. This might leave hardware resources underutilized and may lead to less optimal performance. Out of the maximum of 12 warps per scheduler, this kernel allocates an average of 10.22 active warps per scheduler, but only an average of 0.15 warps were eligible per cycle. Eligible warps are the subset of active warps that are ready to issue their next instruction. Every cycle with no eligible warp results in no instruction being issued and the issue slot remains unused. To increase the number of eligible warps, avoid possible load imbalances due to highly different execution durations per warp. Reducing stalls indicated on the [Warp State Statistics](#) and [Source Counters](#) sections can help, too.

Warps Per Scheduler

Category	Value
GPU Maximum Warps Per Scheduler	12.0
Theoretical Warps Per Scheduler	12.0
Active Warps Per Scheduler	9.0
Eligible Warps Per Scheduler	0.15
Issued Warp Per Scheduler	0.09

Warp State Statistics

Analysis of the states in which all warps spent cycles during the kernel execution. The warp states describe a warp's readiness or inability to issue its next instruction. The warp cycles per instruction define the latency between two consecutive instructions. The higher the value, the more warp parallelism is required to hide this latency. For each warp state, the chart shows the average number of cycles spent in that state per issued instruction. Stalls are not always impacting the overall performance nor are they completely avoidable. Only focus on stall reasons if the schedulers fail to issue every cycle. When executing a kernel with mixed library and user code, these metrics show the combined values.

Warp Cycles Per Issued Instruction [cycle]	110.08	Avg. Active Threads Per Warp	32
Warp Cycles Per Executed Instruction [cycle]	110.11	Avg. Not Predicated Off Threads Per Warp	29.05

long_scoreboard

On average, each warp of this kernel spends 81.9 cycles being stalled waiting for a scoreboard dependency on a L1TEX (local, global, surface, texture, rtcore) operation. This represents about 74.4% of the total average of 110.1 cycles between issuing two instructions. To reduce the number of cycles waiting on L1TEX data accesses verify the memory access patterns are optimal for the target architecture, attempt to increase cache hit rates by increasing data locality or by changing the cache configuration, and consider moving frequently used data to registers and to shared memory.

Warp Stall

Check the [Source Counters](#) section for the top stall locations in your source based on sampling data. The [Kernel Profiling Guide](#) provides more details on each stall reason.

Warp State (All Cycles)

State	Value
Stall Long Scoreboard	81.9
Stall LG Throttle	18.0
Stall Wait	1.0
Selected	0.5
Stall Drain	0.5
Stall MIO Throttle	0.5

Memory Workload Analysis

Detailed analysis of the memory resources of the GPU. Memory can become a limiting factor for the overall kernel performance when fully utilizing the involved hardware units (Mem Busy), exhausting the available communication bandwidth between those units (Max Bandwidth), or by reaching the maximum throughput of issuing memory instructions (Mem Pipes Busy). Detailed chart of the memory units. Detailed tables with data for each memory unit.

Memory Throughput [Gbyte/s]	266.36	Mem Busy [%]	48.62
L1/TEX Hit Rate [%]	25.77	Max Bandwidth [%]	78.48
L2 Hit Rate [%]	7.76	Mem Pipes Busy [%]	78.48
L2 Compression Success Rate [%]	0	L2 Compression Ratio	0

L2 Store Access Pattern

The memory access pattern for stores from L1TEX to L2 is not optimal. The granularity of an L1TEX request to L2 is a 128 byte cache line. That is 4 consecutive 32-byte sectors per L2 request. However, this kernel only accesses an average of 2.0 sectors out of the possible 4 sectors per cache line. Check the [Source Counters](#) section for uncoalesced stores and try to minimize how many cache lines need to be accessed per memory request.

L2 Load Access Pattern

The memory access pattern for loads from L1TEX to L2 is not optimal. The granularity of an L1TEX request to L2 is a 128 byte cache line. That is 4 consecutive 32-byte sectors per L2 request. However, this kernel only accesses an average of 2.0 sectors out of the possible 4 sectors per cache line. Check the [Source Counters](#) section for uncoalesced loads and try to minimize how many cache lines need to be accessed per memory request.

Fig. – Warp Statistics for Matrix Unroll Kernel

Fig. – Memory Workload Analysis for Tiled Matrix Multiplication

d. Does this optimization synergize with any other optimizations? How and why?

Constant memory can be used to store immutable data such as kernel weights that do not change during execution and are shared across threads. When combined with unrolling and tiling, using constant memory reduces memory latency and increases cache efficiency. Furthermore, the mini batch computations can be distributed across different streams to use memory efficiently. Loop unrolling could be implemented to exploit instruction level parallelism in the loops of the unrolling and matrix multiplication. Lastly, this optimization can benefit from kernel fusion which is discussed in the next section.

- e. List your references used while implementing this technique. (you must mention textbook pages at the minimum)
- i. Course Textbook p. 359-366
 - ii. Slide Deck 18: Computation in Deep Neural Networks

4. Optimization Number #3: Kernel Fusion for Unrolling and Matrix-Multiplication

a. How does this optimization work in theory? Expected behavior?

Kernel fusion combines multiple computational steps into a single kernel. This approach reduces the overhead of launching multiple kernels and minimizes data transfers between the device and memory. By fusing the unrolling and matrix multiplication steps from optimization #2, the data is easier to access. Less time is wasted replicating, which significantly speeds up the process by reducing global memory access and improving locality.

b. How did you implement your code? Explain thoroughly and show code snippets.

The fused kernel below performs both unrolling and matrix multiplication using shared memory. This fusion of operations is done within a single kernel launch to ensure data remains on-chip, enhancing computational speed and efficiency. The code snippet below demonstrates how the matrix is unrolled and loaded into shared memory. The comments in the snippet explicitly state how the unrolling occurs. Furthermore, the kernel is configured to handle dimensions appropriately for the input and output sizes, ensuring that the tiling covers the entire operation. 3D grid dimensions are used to accommodate batch processing and tiling within the matrix multiplication.

```

for(int m = 0; m < ceil(numAColumns/(1.0 * TILE_WIDTH)) ; ++m){
    if( Row < numARows && (m * TILE_WIDTH + tx) < numAColumns){
        subTileA[ty][tx] = A[Row * numAColumns + (m * TILE_WIDTH + tx)];
    }
    else{
        subTileA[ty][tx] = 0;
    }

    if((m * TILE_WIDTH + ty) < numBRows && Col < numBColumns){
        i = (m * TILE_WIDTH + ty) * numBColumns + Col;
        rowUnroll = i / W_unroll; //unroll row index
        colUnroll = i % W_unroll; // unroll column index
        c = rowUnroll / (K * K); //calculate channel
        h = colUnroll / Width_out; //height index for output dims
        w = colUnroll % Width_out; // width index for output dims
        p = (rowUnroll/K) % K; //offset in height dim within kernel "window"
        q = rowUnroll % K; //same as p but for width
        subTileB[ty][tx] = in_4d(bz, c, h + p, w + q);
    }
    else{
        subTileB[ty][tx] = 0;
    }

    __syncthreads();

    for(int k = 0; k < TILE_WIDTH; ++k){
        Cvalue += subTileA[ty][k] * subTileB[k][tx];
    }
    __syncthreads();
}

```

Fig. – Code snippet of fused kernel for unrolling and matrix multiplication

- c. Did the performance match your expectation? Show your analysis results using profiling tools.

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
------------	-----------	-----------	----------------------	----------

100	<i>0.637177 ms</i>	<i>0.356009 ms</i>	<i>0m4.044s</i>	<i>0.86</i>
1000	<i>5.81798 ms</i>	<i>3.32418 ms</i>	<i>0m36.556s</i>	<i>0.886</i>
10000	<i>57.4412 ms</i>	<i>33.0359 ms</i>	<i>6m2.886s</i>	<i>0.8714</i>

Compared to optimization #2, this optimization is much faster as it demonstrates the benefits of fusion, but the op times are still slower compared to the baseline. The overall execution time is slightly improved, but this could just vary with each execution of the program. The slow down in op time could be due to how the memory is being accessed and loaded into shared memory causing op times to slow down.

The nsys profiling shows how many fewer kernel launches and cudaDeviceSynchronizes compared to two separate kernels in the previous optimization. With fusion, only one kernel needs to be launched per layer in the convolution. It can also be noticed that fewer time is spent in gpumemtimesum compared to other optimizations.

[5/8] Executing 'cudaapism' stats report										
Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name		
53.4	290,440,778	8	36,305,097.3	10,708,233.5	32,211	146,255,861	55,987,919.9	cudaMemcpy		
28.0	152,305,865	8	19,038,233.1	135,218.0	82,765	151,021,902	53,329,690.4	cudaAlloc		
16.8	91,280,972	8	11,410,121.5	6,372.0	2,044	58,276,648	22,174,521.2	cudaDeviceSynchronize		
1.7	9,432,656	8	1,179,082.0	1,012,375.0	108,803	2,587,724	1,067,867.5	cudaFree		
0.1	290,205	6	48,367.5	36,198.5	15,730	107,421	35,898.2	cudaLaunchKernel		
0.0	1,172	1	1,172.0	1,172.0	1,172	1,172	0.0	cuModuleGetLoadingMode		

[6/8] Executing 'gpukernsum' stats report										
Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ	
63.9	58,272,455	1	58,272,455.0	58,272,455.0	58,272,455	58,272,455	0.0	400	1 10000	16 16 1
36.1	32,970,699	1	32,970,699.0	32,970,699.0	32,970,699	32,970,699	0.0	73	1 10000	16 16 1
0.0	4,864	2	2,432.0	2,432.0	2,400	2,464	45.3	1 1 1	1 1 1	do_not_remove_this_kernel()
0.0	4,672	2	2,336.0	2,336.0	2,304	2,368	45.3	1 1 1	1 1 1	prefn_marker_kernel()

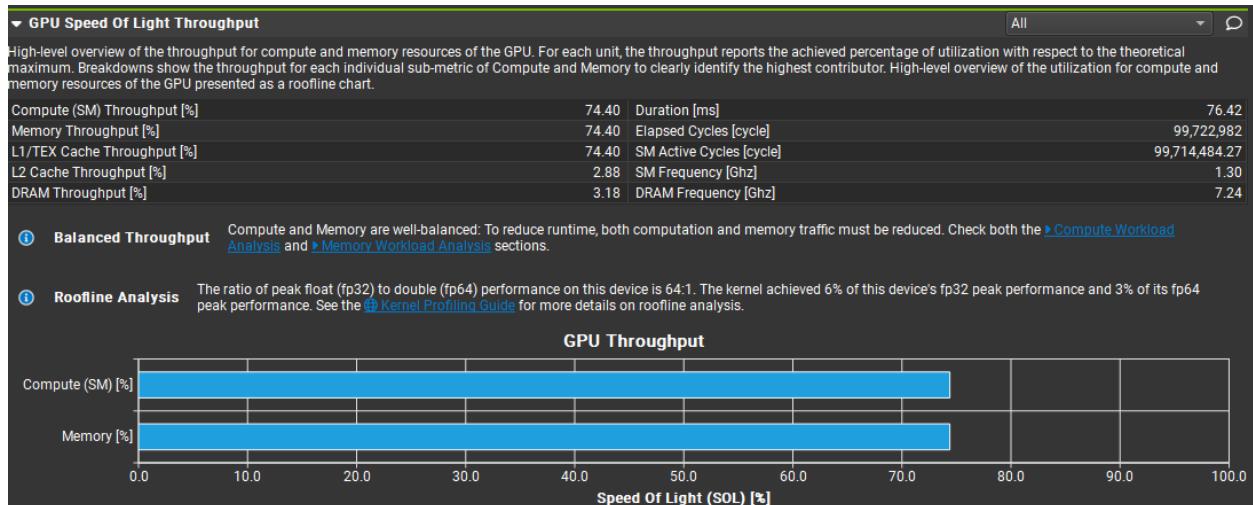
[7/8] Executing 'gpumemtimesum' stats report										
Time (%)	Total Time (ns)	Count	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Operation		
85.5	246,496,761	2	123,203,380.5	123,203,380.5	100,633,223	145,773,538	31,919,022.8	[CUDA memcpy DtoH]		
14.5	41,811,581	6	6,968,596.8	1,823.5	1,472	21,525,775	10,800,335.3	[CUDA memcpy HtoD]		

[8/8] Executing 'gpumemsizesum' stats report							
Total (MB)	Count	Avg (MB)	Med (MB)	Min (MB)	Max (MB)	StdDev (MB)	Operation
1,763.840	2	881.920	881.920	739.840	1,024.000	200.931	[CUDA memcpy DtoH]
551.853	6	91.976	0.007	0.000	295.840	143.039	[CUDA memcpy HtoD]

Fig. – nsys results 5-8 for Kernel Fusion for Unrolling and Matrix-Multiplication

In Nsight Compute, the GPU throughput is well balanced and to reduce runtime, both computation and memory traffic must be reduced as said in the figure below. This throughput is a significant improvement to the separated optimizations. The compute workload analysis demonstrates that FP64 is a heavily utilized pipeline at

73.4%. The LSU pipeline is used heavily to handle low latency instructions. In addition, the L1/TEX and L2 cache are well utilized as they both have a high hit rate, but the L2 store and load access pattern could improve as seen in the memory workload analysis figure. Lastly, the shared memory is utilized shown by the shared memory data, but there still exist 2,896,371 bank conflicts which is slowing down the program. These bank conflicts may be the reason why the op times are still slower



than the baseline.

Fig. – Kernel Fusion for Unrolling and Matrix-Multiplication GPU SOL Throughput

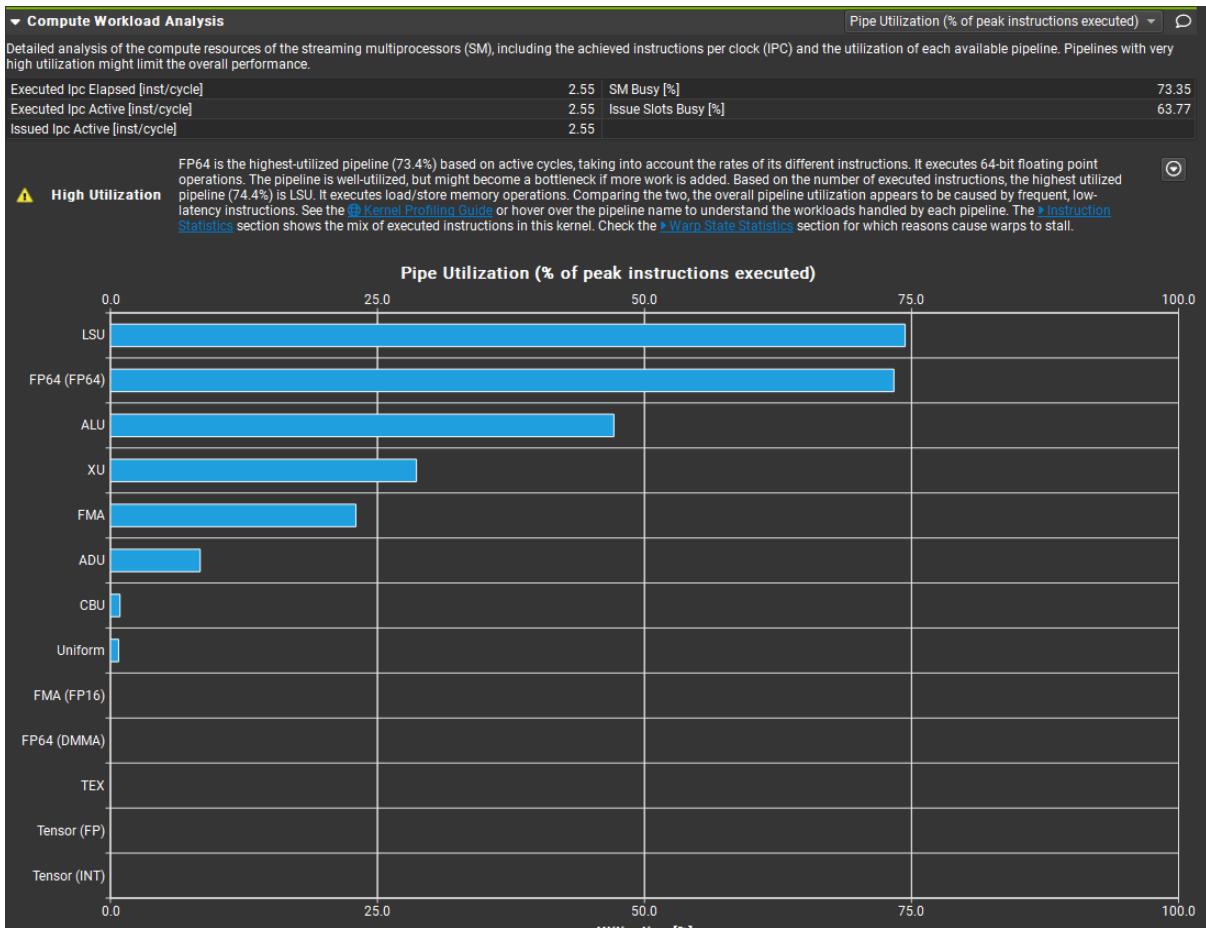


Fig. - Compute Workload Analysis for Kernel Fusion for Unrolling and Matrix-Multiplication

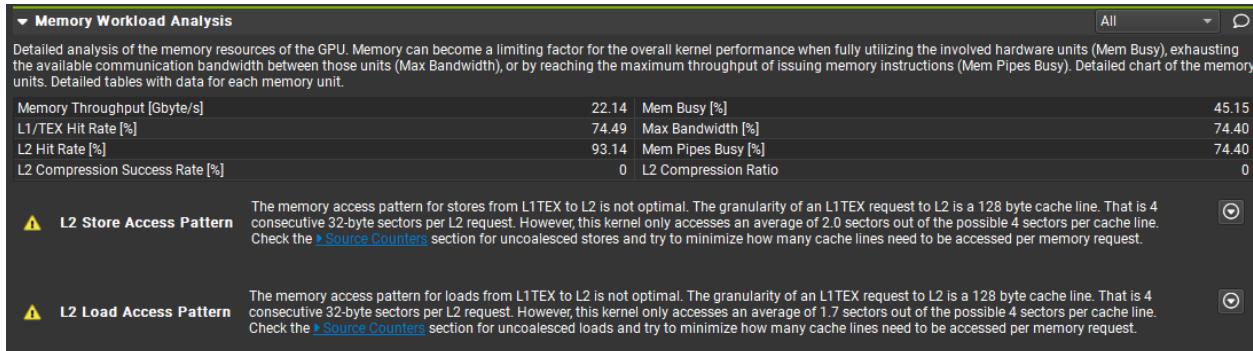


Fig. - Memory Workload Analysis for Kernel Fusion for Unrolling and Matrix-Multiplication

Shared Memory					
	Instructions	Requests	Wavefronts	% Peak	Bank Conflicts
Shared Load	2,560,000,000	2,560,000,000	3,078,385,068	36.75	1,990,252
Shared Load Matrix	0	0			
Shared Store	256,000,000	256,000,000	256,000,000	0.76	0
Shared Store From Global Load	0	0	0	0	0
Shared Atomic	0	0	0	0	0
Other	-	-	193,352,575	4.60	906,119
Total	2,816,000,000	2,816,000,000	3,527,737,643	42.11	2,896,371

Fig. – Shared Memory Data for Kernel Fusion for Unrolling and Matrix-Multiplication

d. Does this optimization synergize with any other optimizations? How and why?

Kernel fusion, when combined with loop unrolling and constant memory usage, could lead to performance improvement. Loop unrolling within the fused kernel reduces the overhead of loop control and increases instruction throughput by executing multiple iterations in a single loop cycle. This enhances the efficiency of the computation within the kernel. Additionally, using constant memory in fused kernels provides fast access to frequently used data (the mask), minimizing memory latency. This synergy improves data throughput and computational speed, as constant data is accessed faster, and loops are executed more efficiently.

- e. List your references used while implementing this technique. (you must mention textbook pages at the minimum)
- i. Course Textbook p. 359-366
 - ii. Slide Deck 18: Computation in Deep Neural Networks

5. Optimization Number #4: Weight Matrix (Kernel) in constant memory

a. How does this optimization work in theory? Expected behavior?

This optimization places the weight matrix in constant memory, which has significantly faster access times than global memory. The weight matrix is placed in constant memory since its values do not change and it is frequently reused. This reduces memory latency, leading to faster convolution operations. By significantly eliminating global memory accesses for the mask, op times can be expected to improve.

b. How did you implement your code? Explain thoroughly and show code snippets.

The kernel weights are placed in constant memory using `__constant__` and accessed via the mask `_4d` macro. The constant memory size calculation comes from the number of inputs and channels there, this can be identified through the output feature map sizes in the kernel launches. The `cudaMemcpyToSymbol` function is used to copy the weights from host memory to constant memory.

```
__constant__ float maskConstant[16*4*7*7];
```

Fig. - Code snippet of constant memory allocation

```
acc += in_4d(n, c, h + p, w + q) * mask_4d(m, c, p, q);
```

Fig. – Code snippet of mask_4d in convolution computation

```
cudaMemcpyToSymbol(maskConstant, host_mask, maskSize * sizeof(float));
```

Fig. – Code snippet of copying mask from host to constant memory

- c. Did the performance match your expectation? Show your analysis results using profiling tools.

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.137679 ms	0.386495 ms	0m4.087s	0.86
1000	1.00568 ms	4.0301 ms	0m36.836s	0.886
10000	10.0485 ms	43.2677 ms	6m3.901s	0.8714

The op times demonstrate that constant memory optimization significantly improves performance. The program's run time improves as there are significantly fewer cycles when accessing constant memory versus global, and with repetitive use of the mask, this is extremely beneficial.

Firstly, in the nsys results below, kernel time has reduced compared to other optimizations and the baseline. Second, there are fewer cudaMemcpys, fewer cudaMallows, and fewer cudaFrees due to the single cudaMemcpyToSymbol. These reductions have led to an increase in performance.

Fig. – nsys results 5-8 for Weight Matrix (Kernel) in Constant Memory

In Nsight Compute, the GPU SOL throughput has decreased, but this can be realized since the kernel is no longer being accessed globally. Even with a decreased throughput, it is still well balanced, and op times have still significantly improved because constant memory is faster. The lower throughput gives space for other optimizations as well. Additionally, looking at the memory workload analysis in the memory chart image, there is a higher memory throughput, 394.99 GB/S, compared to other optimizations that maximize overall throughput. Also, in the memory chart diagram, there is 95.20% hit rate on the L1/TEX Cache (because of the constant memory allocation) which compared to other optimizations is a significant improvement. The yellow icons state that there could be better memory coalescing to take advantage of a memory burst. To improve, the global access pattern would have to change so that the L1/TEX and L2 Cache are utilized effectively.

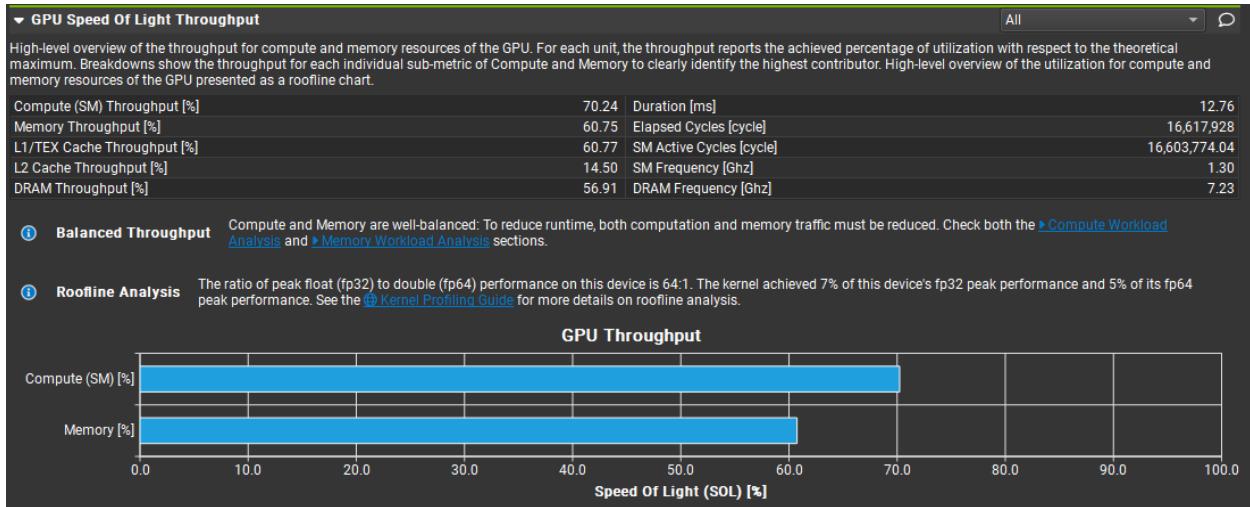


Fig. - Weight Matrix (Kernel) in constant memory GPU SOL Throughput

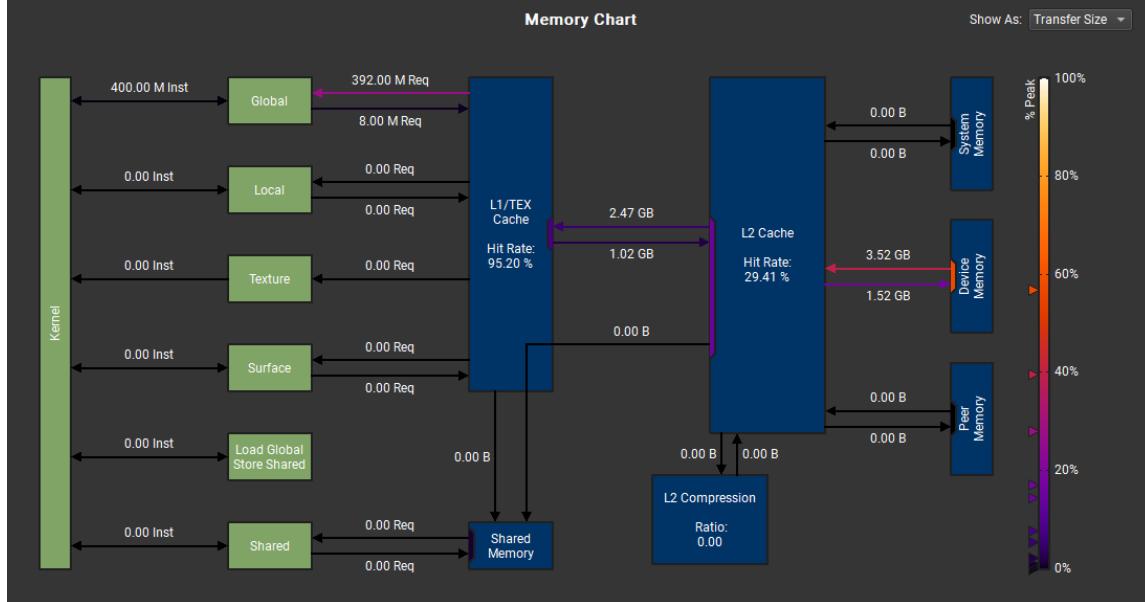
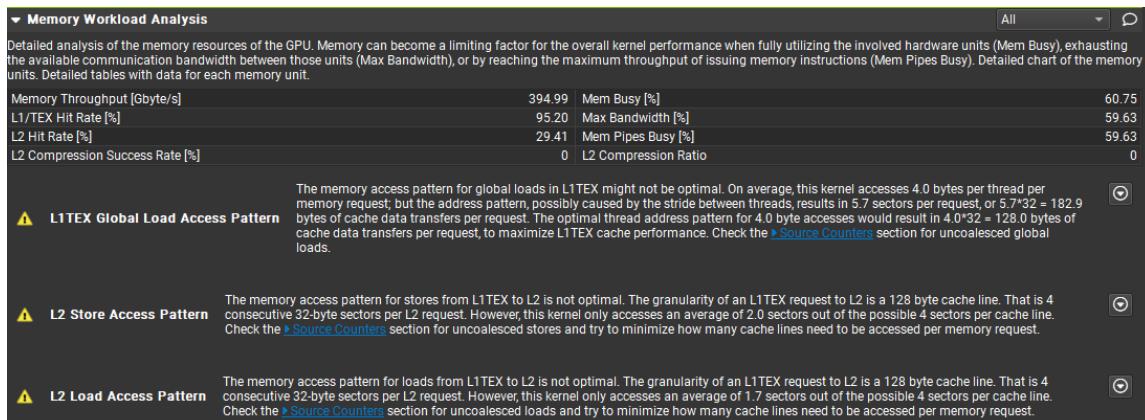


Fig – Nsight Compute Memory Chart for Weight Matrix (Kernel) in constant memory

d. Does this optimization synergize with any other optimizations? How and why?

Yes, it synergizes with optimizations like using shared memory for input tiles, unrolling and matrix multiplication, loop unrolling, kernel fusion, etc. All these optimizations still require the use of the kernel, and since the values in the kernel do not change, using constant memory can reduce global memory accesses and improving data reuse, further enhancing performance. Taking advantage of the fast constant memory for all these optimizations is a clear decision.

e. List your references used while implementing this technique. (you must mention textbook pages at the minimum)

i. Slide Deck 8: Convolution and Constant Memory

6. Optimization Number #5: Tuning with Restrict and Loop Unrolling

a. How does this optimization work in theory? Expected behavior?

For tuning with restrict and loop unrolling, the `__restrict__` keyword informs the compiler that the object pointed to by the pointer will only be accessed by that pointer, ensuring no memory aliasing occurs. This is useful where memory bandwidth and access patterns impact performance. Loop unrolling reduces the overhead of loop control, decreasing conditional branching and increasing instruction throughput. Together, these optimizations are expected to enhance memory access efficiency and computational speed, leading to reduced execution times and increased throughput.

b. How did you implement your code? Explain thoroughly and show code snippets.

The implementation involves marking the input and mask pointers with the `restrict` qualifier to suggest that these pointers will not alias, allowing the compiler to generate more efficient memory access instructions. `__restrict__` is placed in front of the input and mask that are passed in as parameters to take advantage of the read-only cache. Loop unrolling is employed using `#pragma unroll` before loops. The compiler will unroll the loops reducing the need to check the conditional in the loop. `#pragma unroll` is placed before the loops that handle the convolution computation, effectively reducing the number of iterations and the overhead. Below is a code snippet demonstrating the use of `restrict` in the parameters and `#pragma unroll` before the loops.

```
_global__ void conv_forward_kernel(float *output, const float* __restrict__ input, const float* __restrict__ mask,
{
    const int Height_out = Height - K + 1;
    const int Width_out = Width - K + 1;

#define out_4d(i3, i2, i1, i0) output[(i3) * (Map_out * Height_out * Width_out) + (i2) * (Height_out * Width_out * K) + (i1) * (Width_out * K) + i0]
#define in_4d(i3, i2, i1, i0) input[(i3) * (Channel * Height * Width) + (i2) * (Height * Width) + (i1) * (Width * K) + i0]
#define mask_4d(i3, i2, i1, i0) mask[(i3) * (Channel * K * K) + (i2) * (K * K) + (i1) * (K) + i0]

    // Insert your GPU convolution kernel code here
    int W_grid = ceil(Width_out/(1.0 * TILE_WIDTH)); //number of horizontal tiles per output map

    int n = blockIdx.x;
    int m = blockIdx.y;
    int h = (blockIdx.z/W_grid) * TILE_WIDTH + threadIdx.y;
    int w = (blockIdx.z % W_grid) * TILE_WIDTH + threadIdx.x;

    if(w < Width_out && h < Height_out){ //boundary check
        float acc = 0.0;
        #pragma unroll
        for(int c = 0; c < Channel; c++){ //sum over the input channels
            #pragma unroll
            for(int p = 0; p < K; p++){ // loop over kernel
                #pragma unroll
                for(int q = 0; q < K; q++){
                    acc += in_4d(n, c, h + p, w + q) * mask_4d(m, c, p, q); //compute sum
                }
            }
        }
    }
}
```

Fig. – Code snippet of Tuning with `restrict` and loop unrolling

- c. Did the performance match your expectation? Show your analysis results using profiling tools.

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.173955 ms	0.50287 ms	0m4.024s	0.86
1000	1.20073 ms	4.93621 ms	0m36.389s	0.886

10000	11.5396 ms	51.7928 ms	5m59.084s	0.8714
-------	---------------	------------	-----------	--------

Based on the op times and total execution, there is a slight improvement in performance when tuning with restrict and loop unrolling compared to the baseline. This slight improvement was to be expected as the code hadn't changed and only restricts and unrolls were added which add slight performance benefits at a lower level.

The nsys results show that overall kernel times slightly decrease. There is a decrease in gpumemtimesum compared to most optimizations. Memory time is expected to see a decrease as we restrict the inputs and unroll the loops to reduce overhead and utilize the read only cache.

[5/8] Executing 'cudaapisum' stats report									
Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name	
56.8	292,454,507	8	36,556,813.4	10,673,974.5	35,907	146,717,060	56,267,443.9	cudaMemcpy	
29.2	150,266,520	8	18,783,315.0	144,490.0	81,903	148,926,333	52,586,022.1	cudaMalloc	
12.3	63,505,331	6	10,584,221.8	7,904.5	4,930	51,923,493	20,772,237.3	cudaDeviceSynchronize	
1.6	8,423,579	8	1,052,947.4	502,711.0	119,872	2,591,819	1,057,225.7	cudaFree	
0.1	297,798	6	49,633.0	32,160.5	13,556	111,148	37,987.4	cudaLaunchKernel	
0.0	1,172	1	1,172.0	1,172.0	1,172	1,172	0.0	cuModuleGetLoadingMode	

[6/8] Executing 'gpkernsum' stats report									
Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ
81.8	51,920,897	1	51,920,897.0	51,920,897.0	51,920,897	51,920,897	0.0	10000	16 9
18.2	11,552,218	1	11,552,218.0	11,552,218.0	11,552,218	11,552,218	0.0	10000	4 25
0.0	4,832	2	2,416.0	2,416.0	2,400	2,432	22.6	1 1	1 1
0.0	4,704	2	2,352.0	2,352.0	2,336	2,368	22.6	1 1	1 1

[7/8] Executing 'gpumemtimesum' stats report									
Time (%)	Total Time (ns)	Count	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Operation	
85.6	247,854,063	2	123,927,031.5	123,927,031.5	101,782,496	146,071,567	31,317,182.4	[CUDA memcpy DtoH]	
14.4	41,639,756	6	6,939,959.3	1,824.0	1,535	21,988,692	10,774,271.8	[CUDA memcpy HtoD]	

[8/8] Executing 'gpumemsizesum' stats report							
Total (MB)	Count	Avg (MB)	Med (MB)	Min (MB)	Max (MB)	StdDev (MB)	Operation
1,763.840	2	881.920	881.920	739.840	1,024.000	200.931	[CUDA memcpy DtoH]
551.853	6	91.976	0.007	0.000	295.840	143.039	[CUDA memcpy HtoD]

Fig. – nsys results 5-8 for Tuning with Restrict and Loop Unrolling

In the Nsight Compute analysis below, due to the similarity of the code to the baseline, the GPU throughput is very similar and is to be expected. The overall change isn't large, but there is a well balanced high throughput. There could be improvement in the L2 load and store access pattern as seen in the memory workload analysis. The L2 cache has a 29.44% hit rate. Overall, the performance boost for this optimization simply comes from loop unrolling as there are fewer checks for the loop condition and the usage of the read only cache using restrict.

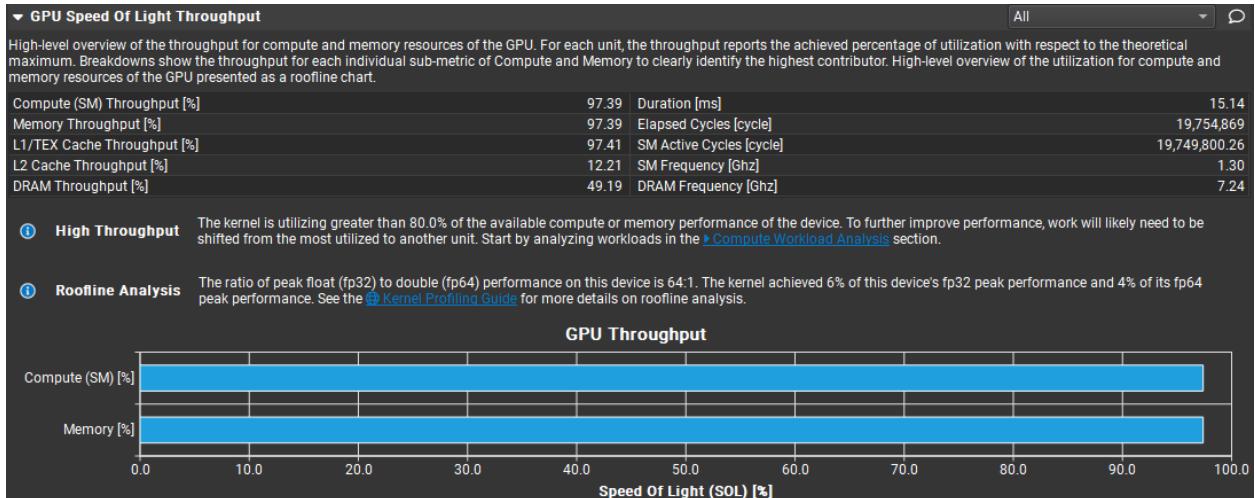


Fig. - Tuning with Restrict and Loop Unrolling GPU SOL Throughput

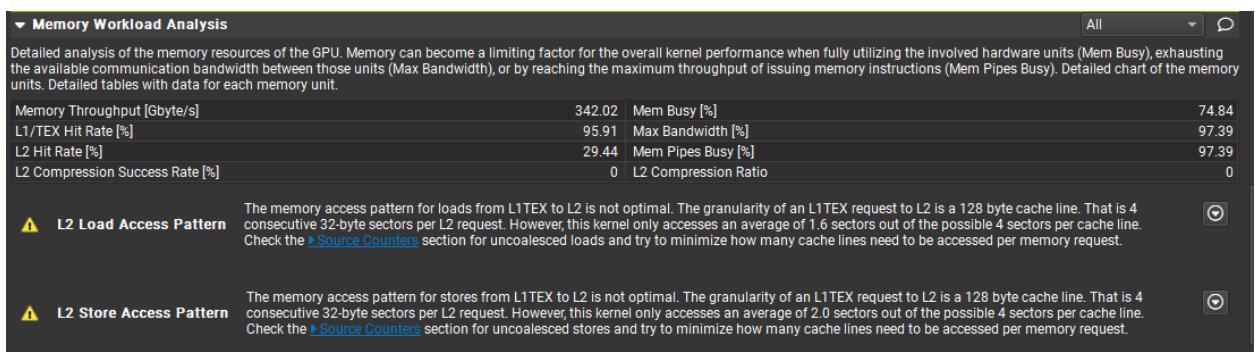


Fig. – Memory Workload Analysis for Tuning with Restrict and Loop Unrolling

d. Does this optimization synergize with any other optimizations? How and why?

This optimization naturally synergizes with all the optimizations in the paper because unrolls could be added in front of most loops and restricts could easily be added to parameters that are only read. Furthermore, shared memory usage becomes more efficient when combined with restrict and loop unrolling, as the predictable access patterns and reduced loop overhead allow for faster and more consistent data handling.

e. List your references used while implementing this technique. (you must mention textbook pages at the minimum)

i. Slide Deck 22: Alternatives to CUDA

- ii. CUDA C++ Programming Guide, Version 12.5, section 7.40 #pragma unroll

7. Optimization Number #8: Input Channel Reduction: Tree

- a. How does this optimization work in theory? Expected behavior?

This optimization uses a reduction tree for collecting computation results across input channels. This method aims to quickly summarize the necessary data before finishing the calculation, making the entire process more efficient. The goal is to take the sum across all channels quickly to produce the output features faster. The expected benefits include a reduction in computational overhead. This optimization aims to minimize global memory accesses and accelerate the overall computation within the convolution process by efficiently summarizing the data before final accumulation.

- b. How did you implement your code? Explain thoroughly and show code snippets.

This code performs a convolution and uses shared memory for efficiency when using the tree. It initializes shared memory, calculates the convolution sum for each thread, and stores it in shared memory. Then, it uses a tree-based reduction to combine these results, halving the contributors in each step until one result per block remains. This result is written to the output by the thread of index zero.

```

if(w < Width_out && h < Height_out){ //boundary check
    float acc = 0.0;
    for(int p = 0; p < K; p++){
        for(int q = 0; q < K; q++){
            acc += in_4d(n, tz, h + p, w + q) * mask_4d(m, tz, p, q);
        }
    }

    reduction[ty * TILE_WIDTH * Channel + tx * Channel + tz] = acc;

    //reduction
    for(unsigned int stride = Channel; stride >= 1; stride >>= 1){
        __syncthreads();

        if(tz < stride && (tz + stride) < Channel){
            reduction[ty * TILE_WIDTH * Channel + tx * Channel + tz] += reduction[ty * TILE_WIDTH * Channel + tx * Channel + tz + stride];
        }
    }

    __syncthreads();

    if(tz == 0){
        out_4d(n, m, h, w) = reduction[ty * TILE_WIDTH * Channel + tx * Channel];
    }
}

```

Fig. – Code snippet for Input Channel Reduction: Tree

- c. Did the performance match your expectation? Show your analysis results using profiling tools.

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.204883 ms	0.871011 ms	0m4.064s	0.86
1000	1.30605 ms	9.79366 ms	0m36.672s	0.886
10000	12.8414 ms	101.913 ms	6m5.675s	0.8714

Despite its theoretical advantages, this optimization technique showed a slowdown in practice compared to the baseline. This could be due to frequent synchronization delays in the reduction steps, inefficient memory access patterns, and underutilization of device resources, especially since the number of channels is 4 in the second layer which isn't very optimal.

Looking at the nsys profiling, it can be noticed that the second layer in the CNN takes up most of the run time, while the first layer is quite fast. The first layer is faster as there is only one input channel so there is no need for a reduction amongst the input channels. The second layer is slow as there are only four input channels per image, so performing a reduction tree with four channels isn't worth it.

[5/8] Executing 'cudaapisum' stats report										
Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name		
54.7	330,359,941	8	41,294,992.6	9,729,248.5	33,252	178,484,876	66,855,986.9	cudaMemcpy		
24.6	148,589,563	8	18,573,695.4	126,582.0	86,452	147,348,139	52,032,979.4	cudaMalloc		
19.2	116,174,501	6	19,362,416.8	7,213.5	4,118	103,367,987	41,470,248.1	cudaDeviceSynchronize		
1.4	8,385,129	8	1,048,141.1	458,102.5	113,212	2,617,711	1,082,318.1	cudaFree		
0.0	239,086	6	39,847.7	29,915.5	16,010	80,431	26,764.7	cudaLaunchKernel		
0.0	1,122	1	1,122.0	1,122.0	1,122	1,122	0.0	cuModuleGetLoadingMode		

[6/8] Executing 'gpukernsum' stats report										
Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ	
89.0	103,363,244	1	103,363,244.0	103,363,244.0	103,363,244	103,363,244	0.0	10000	16	9
11.0	12,778,231	1	12,778,231.0	12,778,231.0	12,778,231	12,778,231	0.0	10000	4	25
0.0	4,800	2	2,400.0	2,400.0	2,368	2,432	45.3	1	1	1
0.0	4,768	2	2,384.0	2,384.0	2,304	2,464	113.1	1	1	1

[7/8] Executing 'gpumemtimesum' stats report									
Time (%)	Total Time (ns)	Count	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Operation	
87.7	288,235,833	2	144,117,916.5	144,117,916.5	110,233,698	178,002,135	47,919,521.4	[CUDA memcpy DtoH]	
12.3	40,293,889	6	6,715,648.2	1,856.0	1,536	21,549,157	10,439,119.6	[CUDA memcpy HtoD]	

[8/8] Executing 'gpumemsizesum' stats report							
Total (MB)	Count	Avg (MB)	Med (MB)	Min (MB)	Max (MB)	StdDev (MB)	Operation
1,763.840	2	881.920	881.920	739.840	1,024.000	200.931	[CUDA memcpy DtoH]
551.853	6	91.976	0.007	0.000	295.840	143.039	[CUDA memcpy HtoD]

Fig. – nsys profiling sections 5-8 for Input Channel Reduction: Tree

In Nsight Compute, it is first noticed that in the first layer, there is a high throughput, while in the second layer the throughput is quite low as there is a latency issue utilizing below 60% throughput. This low throughput could be due to the extremely large number of bank conflicts, 267,002,659, seen in the shared memory data for the second kernel launch (second layer). Additionally, in the second layer, the L2 cache has a low hit rate at 12.44% as shown in the memory workload analysis. It seems that when the reduction is carried out, there are many other complications especially with the use of shared memory and the large count of bank conflicts.

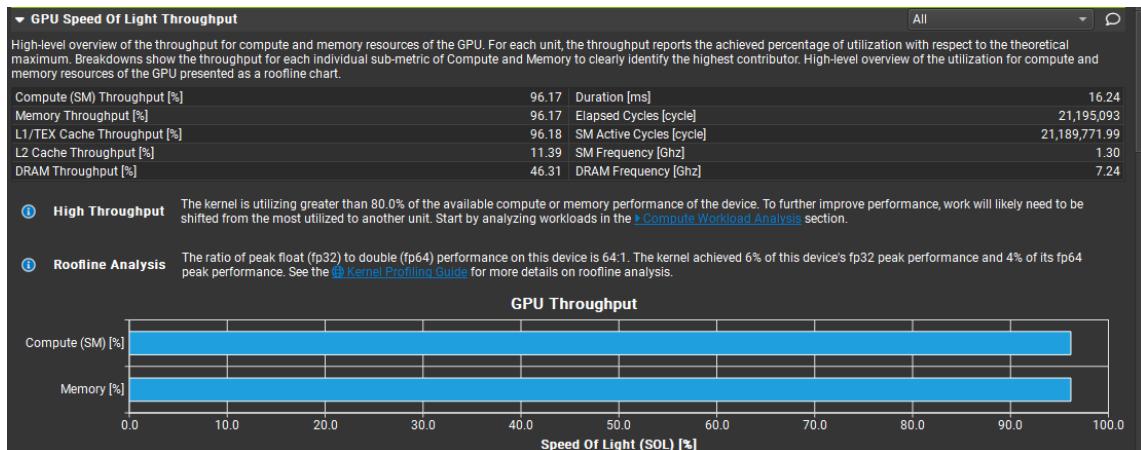


Fig. – GPU SOL Throughput for first layer of Input Channel Reduction: Tree

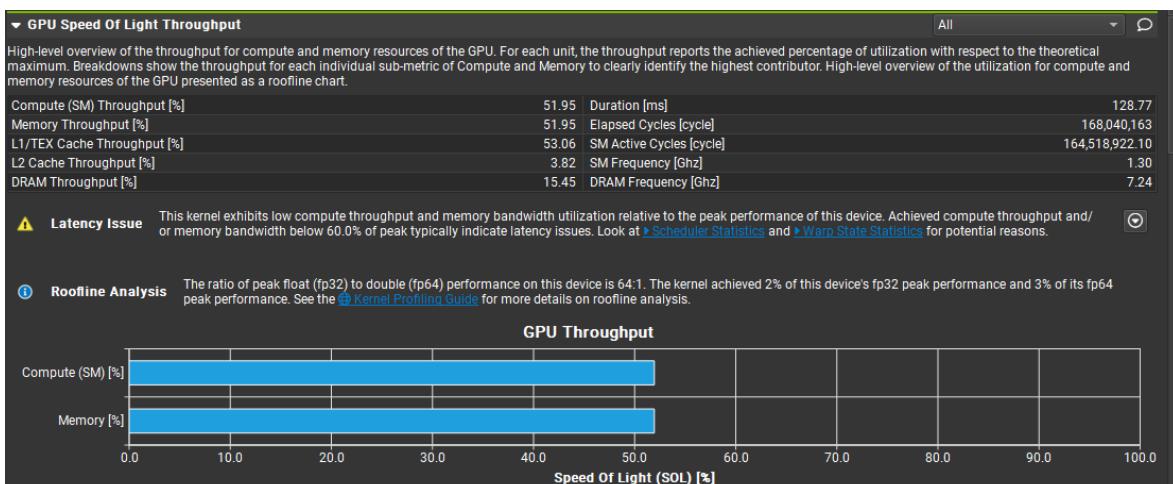


Fig. – GPU SOL Throughput for second layer of Input Channel Reduction: Tree

Shared Memory					
	Instructions	Requests	Wavefronts	% Peak	Bank Conflicts
Shared Load	57,120,000	57,120,000	190,400,000	1.35	133,280,000
Shared Load Matrix	0	0			
Shared Store	57,120,000	57,120,000	190,400,000	0.34	133,280,000
Shared Store From Global Load	0	0	0	0	0
Shared Atomic	0	0	0	0	0
Other	-	-	351,322,659	3.50	442,659
Total	114,240,000	114,240,000	732,122,659	5.19	267,002,659

Fig. – Shared Memory Data for Input Channel Reduction: Tree

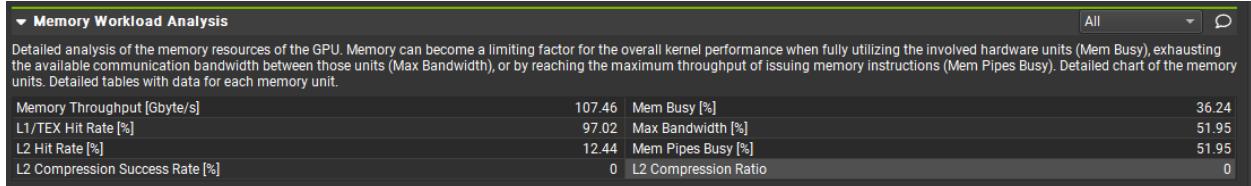


Fig. – Memory Workload Analysis for Input Channel Reduction: Tree

- d. Does this optimization synergize with any other optimizations? How and why?

This reduction optimization complements optimizations such as shared memory optimization, which it already uses, and loop unrolling, which can further minimize loop overhead and maximize arithmetic operations per thread. These enhancements can collectively amplify the computational efficiency and reduce runtime. Additionally, with proper memory coalescing and reducing the number of bank conflicts, if there are more input channels, there could be performance benefits to using a tree.

- e. List your references used while implementing this technique. (you must mention textbook pages at the minimum)
- Slide Deck 11: Reduction Trees

8. Optimization Number #9: *Input Channel Reduction: Atomics*

- a. How does this optimization work in theory? Expected behavior?

The atomic operations optimization is designed to manage data contention effectively when multiple threads simultaneously attempt to write to the same location in memory. This approach is particularly useful for accumulating contributions from multiple channels to the output feature maps without needing a `__syncthreads()`. The expected behavior is a more efficient handling of writes to the output array where each thread

updates the sum independently, reducing the overhead typically associated with thread synchronization.

b. How did you implement your code? Explain thoroughly and show code snippets.

The implementation utilizes atomic addition function (atomicAdd) to safely accumulate the results of the convolution from different threads into the output array. The code is very similar to optimization #8 except instead of using shared memory and a reduction, the reduction is replaced with an atomic add. This is crucial in scenarios where multiple threads compute results that contribute to the same output element, as it prevents race conditions.

```

if(w < Width_out && h < Height_out){ //boundary check
    float acc = 0.0;
    for(int p = 0; p < K; p++){ //loop through the kernel
        for(int q = 0; q < K; q++){
            acc += in_4d(n, tz, h + p, w + q) * mask_4d(m, tz, p, q);
        }
    }
    atomicAdd(&out_4d(n, m, h, w), acc);
}

```

Fig. – Code snippet of Input Channel Reduction: Atomics

c. Did the performance match your expectation? Show your analysis results using profiling tools.

Batch Size	Op Time 1	Op Time 2	Total Execution Time	Accuracy
100	0.177893 ms	0.789646 ms	0m4.011s	0.86
1000	1.23126 ms	8.65766 ms	0m36.655s	0.886
10000	11.9356 ms	92.0994 ms	6m3.722s	0.8714

The use of atomic operations slows down overall op times. There is an increase in time in the second layer as there is a greater number of channels resulting in a larger number of atomic operations needed. The first layer only has one channel so there is

only one atomic operation per output. This could be due to increased memory latency and the serialization of accesses.

Viewing the nsys profiling below, there aren't many differences apart from the op times for the kernel, and the overall gpumemtimesum. The gpumemtimesumhas has decreased compared to other optimizations, and this could be due to fewer memory operations in separate location and the reduction of intermediate storage.

[5/8] Executing 'cudaapisum' stats report									
Time (%)	Total Time (ns)	Num Calls	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name	
51.8	295,095,781	8	36,886,972.6	10,607,130.0	37,450	148,787,588	57,030,106.0	cudaMemcpy	
28.2	160,799,073	8	20,099,884.1	134,326.5	79,960	159,490,472	56,322,547.0	cudaMalloc	
18.3	104,452,758	6	17,408,793.0	7,369.0	3,887	92,447,906	37,072,199.4	cudaDeviceSynchronize	
1.7	9,421,928	8	1,177,741.0	971,785.5	109,025	2,661,414	1,079,468.5	cudaFree	
0.1	324,006	6	54,001.0	41,633.5	17,322	118,822	37,817.5	cudaLaunchKernel	
0.0	1,042	1	1,042.0	1,042.0	1,042	1,042	0.0	cuModuleGetLoadingMode	

[6/8] Executing 'gpukernsum' stats report									
Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	GridXYZ	BlockXYZ
88.5	92,442,501	1	92,442,501.0	92,442,501.0	92,442,501	92,442,501	0.0	10000	16 9 16 16 4
11.5	11,977,340	1	11,977,340.0	11,977,340.0	11,977,340	11,977,340	0.0	10000	4 25 16 16 1
0.0	4,833	2	2,416.5	2,416.5	2,400	2,433	23.3	1 1 1 1 1	do_not_remove_this_kernel()
0.0	4,705	2	2,352.5	2,352.5	2,305	2,400	67.2	1 1 1 1 1	prefn_marker_kernel()

[7/8] Executing 'gpumemtimesum' stats report									
Time (%)	Total Time (ns)	Count	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Operation	
85.8	250,931,679	2	125,465,839.5	125,465,839.5	102,700,002	148,231,677	32,195,756.2	[CUDA memcpy DtoH]	
14.2	41,495,246	6	6,915,874.3	1,824.0	1,504	21,755,795	10,730,535.1	[CUDA memcpy HtoD]	

[8/8] Executing 'gpumemsizesum' stats report							
Total (MB)	Count	Avg (MB)	Med (MB)	Min (MB)	Max (MB)	StdDev (MB)	Operation
1,763.840	2	881.920	881.920	739.840	1,024.000	200.931	[CUDA memcpy DtoH]
551.853	6	91.976	0.007	0.000	295.840	143.039	[CUDA memcpy HtoD]

Fig. – nsys profiling for Input Channel Reduction: Atomics

In Nsight Compute, the first layer has a high throughput while the second layer's throughput is nearly halved as shown below. These results are very similar to the previous optimization as they operate in a similar manner. While the first layer functions well, the second layer has several issues. As shown by the yellow icon in the throughput, there is a latency issue which could be due to the slower atomic operations. The L2 load access pattern is not optimal as well as the code could have a better use of memory coalescing. Additionally, in the scheduler statistics next to the yellow icon, it is shown that the scheduler only issues an instruction every 2.2 cycles, underutilizing device resources. It also mentions how there is a maximum of 12 warps per scheduler, but the kernel allocates 7.13 warps per scheduler, and only 1.11 warps were eligible within that cycle. Lastly, in the warp state statistics, thread divergence is shown to be an issue as the kernel only achieves 23.5 active threads per cycle. With many issues resulting in the use

of atomic operations, atomic operations must only be used when absolutely needed and very carefully.

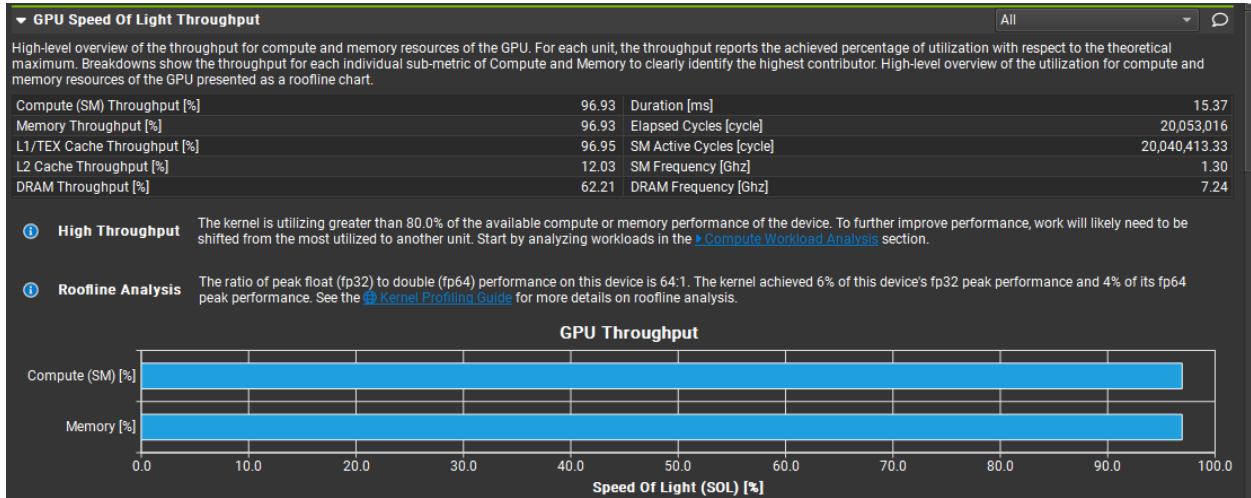


Fig. – GPU SOL Throughput for first layer of Input Channel Reduction: Atomics

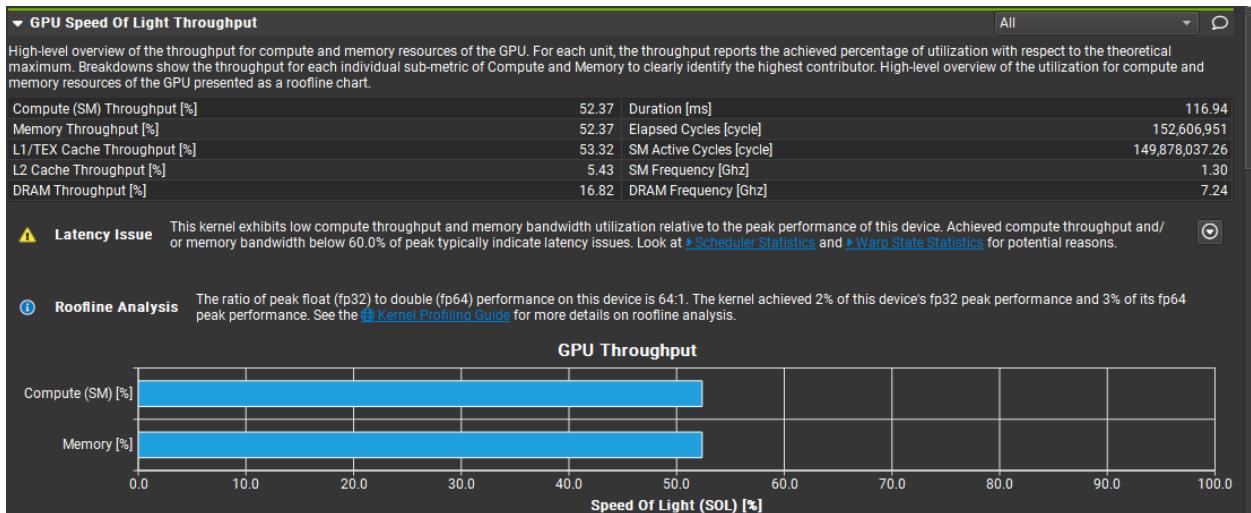


Fig. – GPU SOL Throughput for second layer of Input Channel Reduction: Atomics



Fig. – Scheduler Statistics for second layer of Input Channel Reduction: Atomics

Image on next page

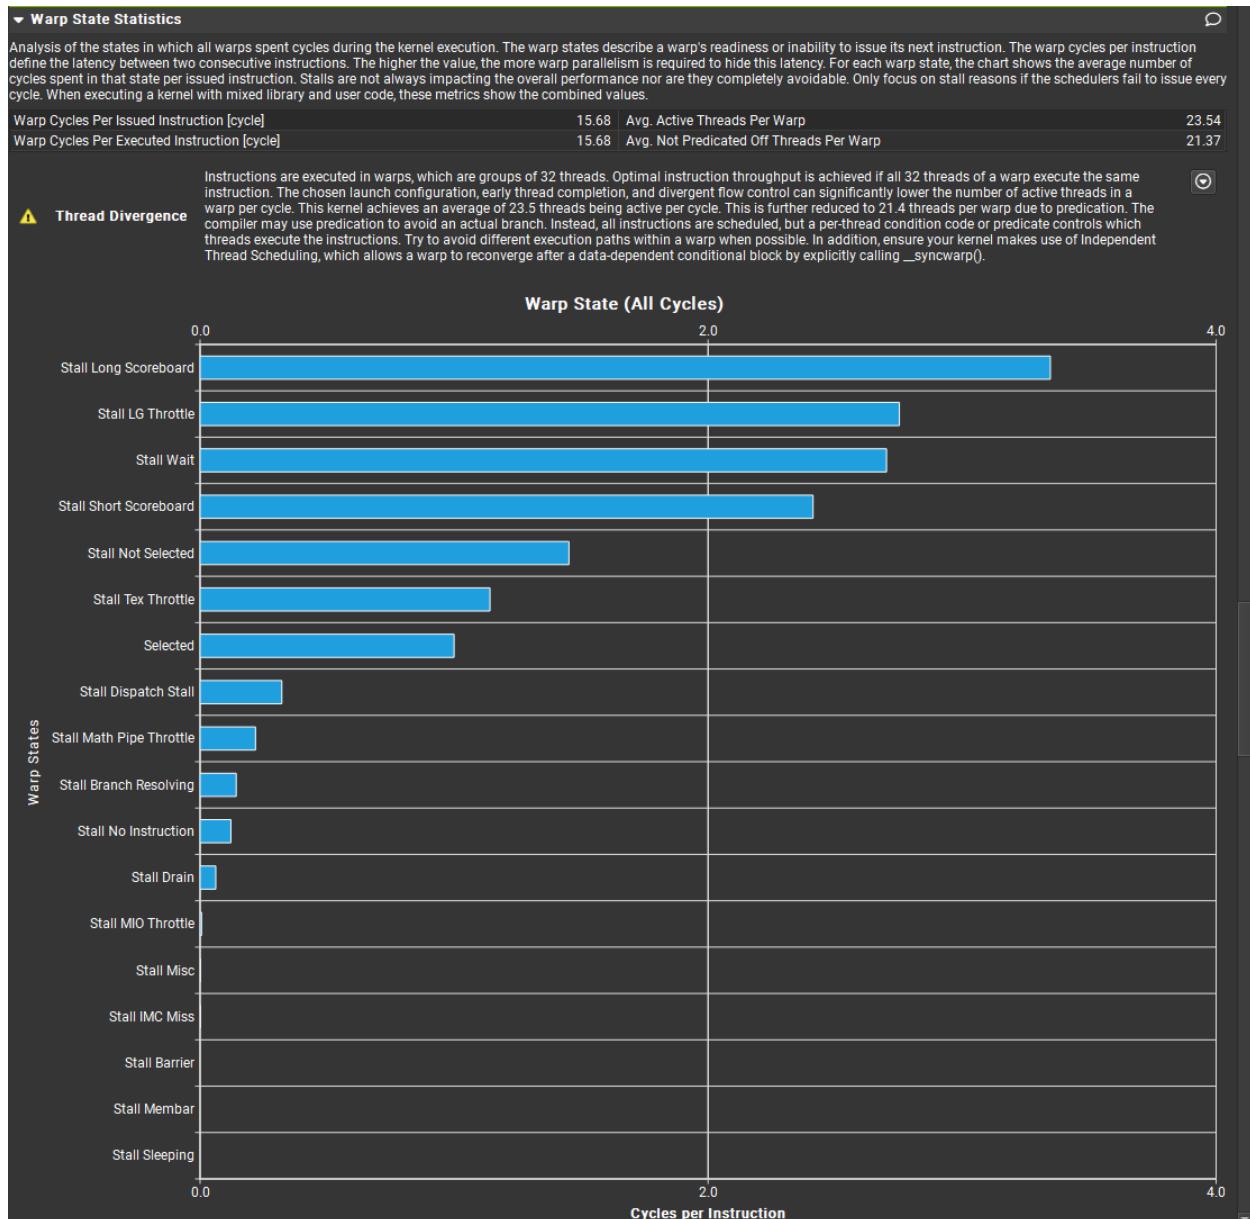


Fig. – Warp State Statistics for second layer of Input Channel Reduction: Atomics

d. Does this optimization synergize with any other optimizations? How and why?

Yes, this optimization synergizes well with optimizations that involve segmenting work among multiple threads, such as tiling or using shared memory. By reducing the granularity of the data each thread works on and then combining results via atomics, the overall memory coherence and execution efficiency can be enhanced. It also pairs well with optimizations that aim to balance the load among threads, ensuring that the atomic operations do not become a bottleneck. Since many of the optimizations such as tiling

and tuning with restrict and loop unrolling require threads to write to output features simultaneously, atomic operations could be used. It is important to use atomic operations only when needed as they have many drawbacks.

List your references used while implementing this technique. (you must mention textbook pages at the minimum)

- i. Slide Deck 13: Atomic Operations and Histograms